

Topics for Module 1

- Part 1: Course Logistics

- **Part 2: Review number representation**

- Fixed point numbers
 - → Floating point numbers

- Textbook Section 3.5 (up to “Floating Point Addition) and the Elaboration part at the end of the section
 - Worksheet on floating point number representation

Quote of the day

**“95% of the
folks out there are
completely clueless about
floating-point.”**

**James Gosling
Sun Fellow
Java Inventor
1998-02-28**



Review of Numbers

- Computers are made to deal with numbers
- What can we represent in N bits?
 - Unsigned integers:

0 to $2^N - 1$

■ Signed Integers (Two's Complement)

$-2^{(N-1)}$ to $2^{(N-1)} - 1$

Other Numbers

• What about other numbers?

- Very large numbers? (seconds/century)
 $3,155,760,000_{10}$ ($3.15576_{10} \times 10^9$)

- Very small numbers? (atomic diameter)
 0.00000001_{10} ($1.0_{10} \times 10^{-8}$)


- Rationals (repeating pattern)
 $2/3$ ($0.666666666\ldots$)

- Irrationals
 $2^{1/2}$ ($1.414213562373\ldots$)

- Transcendentals
 e ($2.718\ldots$), π ($3.141\ldots$)

• All represented in scientific notation

Scientific Notation (in Decimal)



● Scientific Notation:

- exactly one digit to the left of decimal point

● Normalized scientific notation : no leadings 0s

● Alternatives to representing 1/1,000,000,000

- Normalized: 1.0×10^{-9}

- Not normalized: $0.1 \times 10^{-8}, 10.0 \times 10^{-10}$

Scientific Notation (in Binary)

$1.0_{\text{two}} \times 2^{-1}$

The diagram illustrates the components of binary scientific notation. The number $1.0_{\text{two}} \times 2^{-1}$ is shown. An arrow points from the label "exponent" to the power of 2, 2^{-1} . Another arrow points from the label "radix (base)" to the base 2, 2 . A third arrow points from the label "binary point" to the decimal point in the number 1.0_{two} .

• Floating point

- Computer arithmetic that represents numbers where binary point is not fixed
- Declare such variable in C as `float`

“Father” of the Floating point standard

**IEEE Standard 754
for Binary Floating-
Point Arithmetic.**




**1989
ACM Turing
Award Winner!**

Prof. Kahan

www.cs.berkeley.edu/~wkahan/ieee754status/

Single-Precision Floating Point

- Normal format: $\pm 1.\text{xxxxxxxxxx}_{\text{two}} * 2^{\text{yyyy}_{\text{two}}}$



- Sign and Magnitude Representation (32 bits)

- S: Sign (1 for negative, 0 for positive)
- Exponent: yyyy
- Fraction: xxxx...xxxx (Significand: $1 + 0.xxxxx...xxxx$)
=> 24 bits

- Real value = $(-1)^S \times (1+F) \times 2^{(E-127)}$

Why this Order of Fields

- Sort records with FP numbers using integer compares
- Break FP number into 3 parts
 - compare signs: negative < positive
 - compare exponents: big exponent => bigger number
 - compare significands for same exponents

Biased Exponent

- **Bias:** the number subtracted from the normal and unsigned representation to get the real value
- **IEEE 754 bias = $2^{8-1}-1=127$ for single precision**
 - Stored exponent = real exponent + bias
- **$0 < E < 255$**
- **Magnitude of numbers**
 - 2^{-126} to 2^{127}
- **Why biased?**
 - negative exponents appear big in 2's complements

Negative Exponent

- 2's compliment? 1.0×2^{-1} vs. $1.0 \times 2^{+1}$ ($1/2$ vs. 2)

1/2	0	1111 1111	000 0000 0000 0000 0000 0000 0000
2	0	0000 0001	000 0000 0000 0000 0000 0000 0000

- $1/2$ seems bigger than 2

- Desirable notation for exponent

- 0000 0000 : smallest
- 1111 1111 : largest

- Biased exponent:

- -1: 126; 1: 128

1/2	0	0111 1110	000 0000 0000 0000 0000 0000 0000
2	0	1000 0000	000 0000 0000 0000 0000 0000 0000

Single-Precision Range

- **Exponents 00000000 and 11111111 reserved**

- **Smallest value**

- Exponent: 00000001

- $\Rightarrow \text{actual exponent} = 1 - 127 = -126$

- Fraction: 000...00 \Rightarrow significand = 1.0

- $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$

- **Largest value**

- exponent: 11111110

- $\Rightarrow \text{actual exponent} = 254 - 127 = +127$

- Fraction: 111...11 \Rightarrow significand ≈ 2.0 (see next slide)

- $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

Review: Sum of Geometric Series

- A geometric sequence/series: $\{a, ar, ar^2, ar^3, \dots\}$

where:

- a is the first term, and
- r is the factor between the terms (called the "common ratio")

- Sum of a geometric series: $a + ar + ar^2 + \dots + ar^{(n-1)}$

$$\sum_{k=0}^{n-1} (ar^k) = a \left(\frac{1 - r^n}{1 - r} \right)$$

- $0.\overline{1}_{\text{two}} = 2^{-1} + 2^{-2} + \dots + 2^{-n}$, where: $a = 2^{-1}$, $r = 2^{-1}$
 $= (2^{-1})((1 - (2^{-1})^n) / (1 - 2^{-1})) = 1 - 2^{-n}$

Overflow and Underflow

● **Overflow**

- result too large ($> 3.4 \times 10^{38}$)
- Exponent larger than represented in 8-bit Exponent field


● **Underflow**

- result too small? ($> 0, < 1.2 \times 10^{-38}$)
- Negative exponent larger than represented in 8-bit Exponent field

● **How to reduce chances of overflow or underflow?**

Double Precision Fl. Pt. Representation

● **Multiple Word Size (64 bits)**



Fraction (cont'd)

32 bits

● **Double Precision (vs. Single Precision)**

- C variable declared as `double`
- Bias= $2^{11-1}-1=1023$
- $[2.0 \times 10^{-308}, 2.0 \times 10^{308}]$
- greater accuracy due to larger significand

Double-Precision Range

- **Exponents 0000...00 and 1111...11 reserved**

- **Smallest value**

- Exponent: 00000000001
⇒ actual exponent = $1 - 1023 = -1022$
- Fraction: 000...00 ⇒ significand = 1.0
- $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$

- **Largest value**

- Exponent: 11111111110
⇒ actual exponent = $2046 - 1023 = +1023$
- Fraction: 111...11 ⇒ significand ≈ 2.0
- $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

Floating-Point Precision

● **Relative precision**

- all fraction bits are significant
- Single: approx 2^{23}
 - Equivalent to $23 \times \log_{10}2 \approx 23 \times 0.3 \approx 7$ decimal digits of precision
- Double: approx 2^{52}
 - Equivalent to $52 \times \log_{10}2 \approx 52 \times 0.3 \approx 16$ decimal digits of precision

Single Precision vs. Double Precision

	Single Precision	Double Precision
C programming	<code>float, real</code>	<code>double</code>
Size	32 bits, 1 word	64 bits, 2 words
Exponent	8 bits	11 bits
Fraction	23 bits	52 bits
Bias	127	1023
Range	10^{-38} to 10^{38}	10^{-308} to 10^{308}

- **Double Precision has greater range and accuracy**

- Fit all 32-bit in fraction

Understanding the Significand

● Method 1 (Fractions):

- In decimal: $0.340_{10} \Rightarrow 340_{10}/1000_{10}$
 $\Rightarrow 34_{10}/100_{10}$
- In binary: $0.110_2 \Rightarrow 110_2/1000_2 = 6_{10}/8_{10}$
 $\Rightarrow 11_2/100_2 = 3_{10}/4_{10}$

● Method 2 (Place Values):

- Convert from scientific notation
- In decimal: $1.6732 = (1 \times 10^0) + (6 \times 10^{-1}) + (7 \times 10^{-2}) + (3 \times 10^{-3}) + (2 \times 10^{-4})$
- In binary: $1.1001 = (1 \times 2^0) + (1 \times 2^{-1}) + (0 \times 2^{-2}) + (0 \times 2^{-3}) + (1 \times 2^{-4})$
- Interpretation of value in each position extends beyond the decimal/binary point

Special Numbers

- What have we defined so far for single precision floating point numbers?

Exponent	Fraction	Object
0	0	???
0	Nonzero	???
1-254	Anything	+/- fl. Pt #
255	0	???
255	Nonzero	???

Representation for $\pm \infty$

- In FP, divide by 0 should produce $\pm \infty$, not overflow.

- Why?

- OK to do further computations with ∞
 - e.g., $X/0 > Y$ may be a valid comparison

- Representation

- Exponent: 255 (for S.P), 2047 (for D.P)
 - Fraction: 0

Representation for 0

● Represent 0?

- Exponent: 0

- Fraction: 0

- What about sign?

- +0: 0 00000000 0000000000000000000000000000

- -0: 1 00000000 0000000000000000000000000000

● Why two zeroes?

- Helps in some limit comparisons

Representation for Not a Number

• What is $\sqrt{-4}$ or $0/0$?

- If ∞ not an error, these shouldn't be either.
- Called Not a Number (**NaN**)
- Exponent: 255, Fraction: nonzero

• Why is this useful?

- help with debugging
- $\text{op}(\text{NaN}, X) = \text{NaN}$

Representation for Denorms (1/2)


● Problem: a gap among representable FP numbers around 0

- Smallest representable positive num:

$$a = 1.0\ldots_2 \cdot 2^{-126} = 2^{-126}$$

- $a - 0 = 2^{-126}$


Normalization
and implicit 1
is to blame!



Representation for Denorms (2/2)

• Solution for Denormalized numbers

- Gradual underflow
- Representation: Exponent = 0, Fraction=nonzero
- Interpretation: no leading 1 before binary point,
implicit exponent = -126.
- Smallest representable positive num: 2^{-149}
- Second smallest representable positive num: 2^{-148}



Normalized vs. Denormal Numbers

<https://www3.ntu.edu.sg/home/ehchua/programming/java/datarpresentation.html>

Not all real numbers
in the range are representable

$-N_{\max}$

$-N_{\min}$

$+N_{\min}$

$+N_{\max}$

Normalized floating-point numbers

0

$-D_{\max}$


$-D_{\min}$

$+D_{\min}$

$+D_{\max}$

Denormalized floating-point numbers

IEEE 754 Floating-Point Numbers



Single Precision		Double Precision		Object Represented
Exponent	Fraction	Exponent	Fraction	
0	0	0	0	0
0	Non-zero	0	Nonzero	\pm Denormalized number
1 - 254	Anything	1 - 2046	Anything	\pm FP number
255	0	2047	0	\pm infinity
255	Nonzero	2047	Nonzero	NaN

Converting Binary Fractions to Decimal FP

0	0110 1000	101 0000 0000 0000 0000 0010
---	-----------	------------------------------

- **Sign:** 0 => positive

- **Exponent:**

- $0110\ 1000_{\text{two}} = 104_{\text{ten}}$

- Bias adjustment: $104 - 127 = -23$

- **Significand:**

- $1 + 1 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3} + 1 \times 2^{-22} + \dots$
 $= 1 + 2^{-1} + 2^{-3} + 2^{-22}$

- **Final answer:**

- $(1+2^{-1}+2^{-3}+2^{-22}) * 2^{-23}$

Converting Decimal Fractions to Binary

● Simple Case

- denominator is power of 2 (2, 4, 8, 16, etc.)

● Show F.P. representation of -0.75

- $-0.75 = -3/4$

- $-11_{\text{two}}/100_{\text{two}} = -0.11_{\text{two}}$

- Normalized to $-1.1_{\text{two}} \times 2^{-1}$

- $(-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent}-127)}$

- $(-1)^1 \times (1 + .100\ 0000\ ... \ 0000) \times 2^{(126-127)}$

1	0111 1110	100 0000 0000 0000 0000 0000
---	-----------	------------------------------

Converting Decimal Fractions to Binary

● Complicated case

- Denominator is not power of 2
- Cannot represent the number precisely

● How to get the significand of a never-ending number?

■ Fact:

- All rational numbers have a repeating pattern in decimal and binary

■ For conversion:

- Write out binary number with repeating pattern.
- Cut it off after correct number of bits
 - different for single vs. double precision
- Derive Sign, Exponent and Significand fields

Converting Decimal to Binary FP: Method1

• 1/3

$$= 0.3333\dots_{10}$$

$$= 0.25 + 0.0625 + 0.015625 + 0.00390625 + \dots$$

$$= 1/4 + 1/16 + 1/64 + 1/256 + \dots$$

$$= 2^{-2} + 2^{-4} + 2^{-6} + 2^{-8} + \dots$$

$$= 0.0101010101\dots_2 * 2^0$$

$$= 1.0101010101\dots_2 * 2^{-2}$$

■ Sign: 0

■ Exponent = $-2 + 127 = 125 = 01111101$

■ Fraction = 0101010101...

0	0111 1101	0101 0101 0101 0101 0101 010
---	-----------	------------------------------

Converting Decimal Fractions to Binary FP: Method 2

- **Use a table with columns old, bit, and new**
- **Work across each row**
 - Multiply the old entry by 2 to get a new value
 - Put the integer part of the new value in the bit entry
 - Put the fraction part in the new entry
- **If the new entry is non-zero, copy it to the old entry in the next row down and continue; o/w, done.**
- **Example**

old	bit	new
.375	0	.75
0.75	1	.5
.5	1	0

So the result is .011

In-class Exercises

- Convert 20_{ten} to single-precision floating point format
- Convert 0.1_{ten} to single-precision floating point format

Answer

- $20_{\text{ten}} = 10100_{\text{two}} = 1.01 * 2^4 = 0\ 10000011\ 01000000\dots$
 - Biased exponent = $4+127 = 131$

$$\begin{aligned}0.1_{\text{ten}} &= 0.0\ 0011\ 0011\ 0011\dots_{\text{two}} \\&= 1.1\ 0011\ 0011\ 0011\dots * 2^{-4}\end{aligned}$$

0 01111011 1 0011 0011 0011 0011 0011 00