

Topics

- Single Cycle CPU Design (sections 4.1-4.4)

- **Pipelining**

- → Overview (section 4.6)
- Pipelined Datapath and control (section 4.7)
- Data hazard (section 4.8)
- Control hazard (section 4.9)

***Thus times do shift,
each thing his turn does hold;
New things succeed,
as former things grow old.***

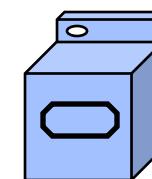
Robert Herrick

Pipelining is Natural!

- Pipelining is an implementation technique in which multiple instructions are overlapped in execution

- Laundry Example

- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, and fold
 - Washer takes 30 minutes
 - Dryer takes 40 minutes
 - “Folder” takes 20 minutes



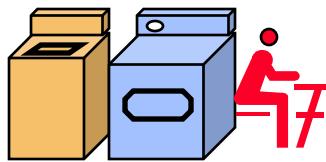
Sequential Laundry

6 PM 7 8 9 10 11 Midnight

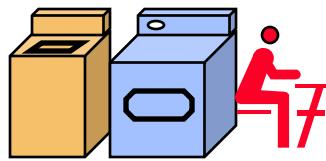


Task Orderer

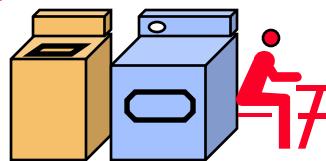
A



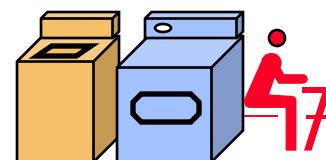
B



C

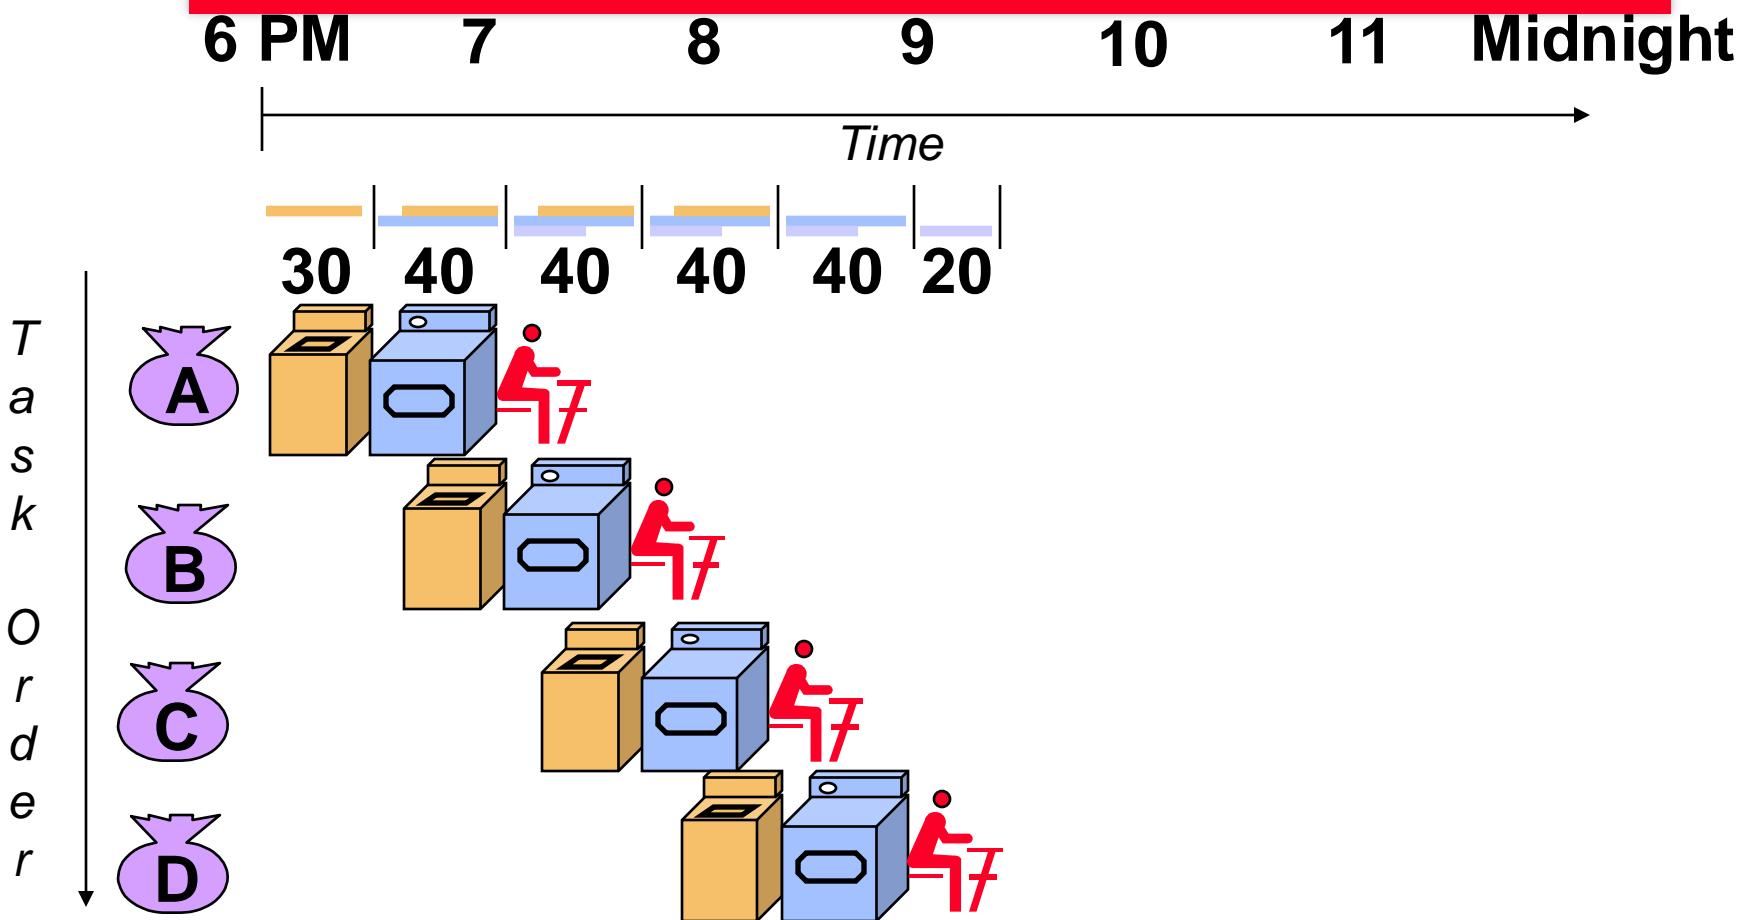


D



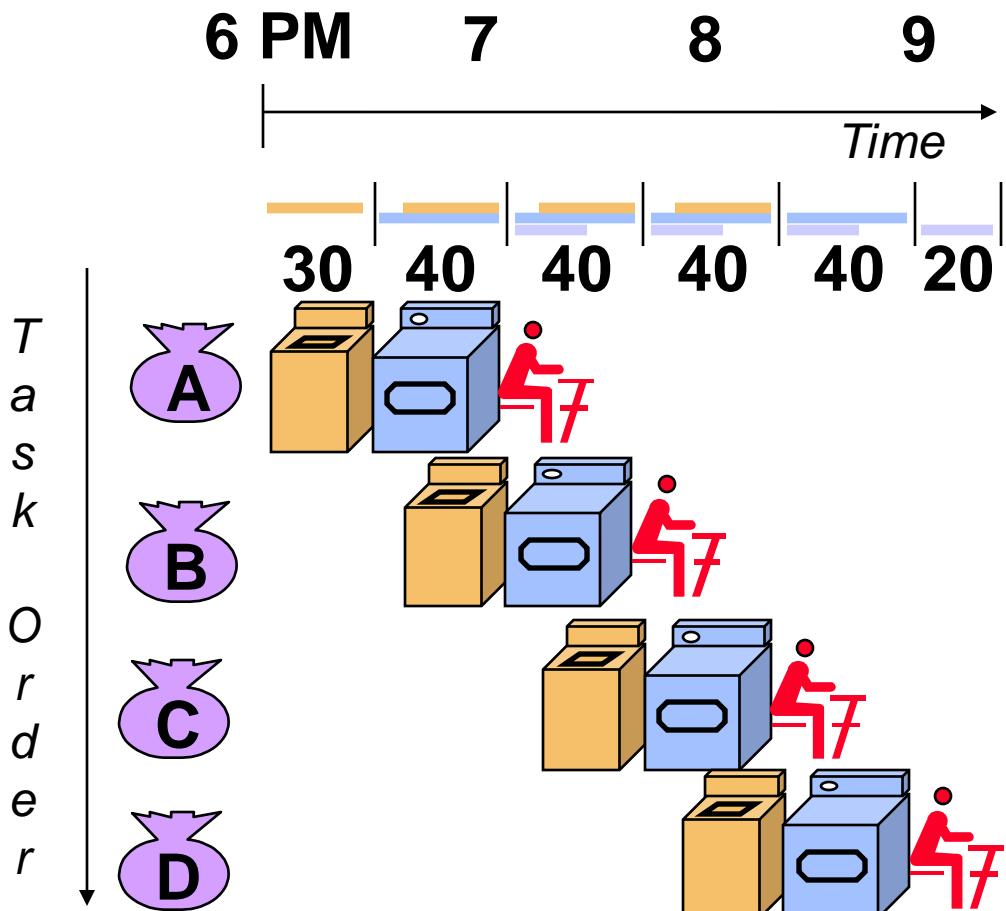
- Sequential laundry takes 6 hours for 4 loads
- If they learned pipelining, how long would laundry take?

Pipelined Laundry: Start work ASAP



- Pipelined laundry takes 3.5 hours for 4 loads
- $30 + 3 \times 40 + (40+20) = 210$ minutes

Pipelining Lessons

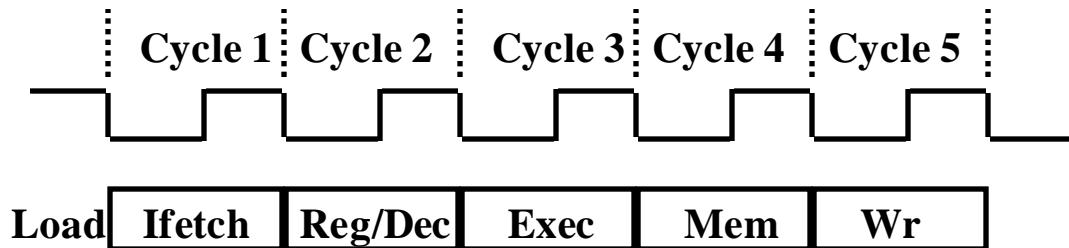


- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- Pipeline rate is limited by **slowest** pipeline stage
- Unbalanced lengths of pipeline stages reduces speedup
- Multiple tasks operating simultaneously using different resources
- Potential speedup = the number of pipeline stages
- Time to “fill” pipeline and time to “drain” it reduces speedup
- Stall for Dependencies

Ideal Speedup

- **Ideal**
 - Each stage takes the same amount of time (t)
 - Each instruction needs the same number of stages (s)
 - The number of instructions (n) is infinite ∞
- $ET_{\text{sequential}} = nst$
- $ET_{\text{pipelined}} = st + (n-1)t = (n-1+s)t$
- $\text{Speedup} = ET_{\text{sequential}}/ET_{\text{pipelined}} = ns/(n-1+s)$
- **Ideal speedup**
= $\lim_{n \rightarrow \infty} \text{speedup}$
= $\lim_{n \rightarrow \infty} (ns/(n-1+s))$
= s

The Five Stages of LW

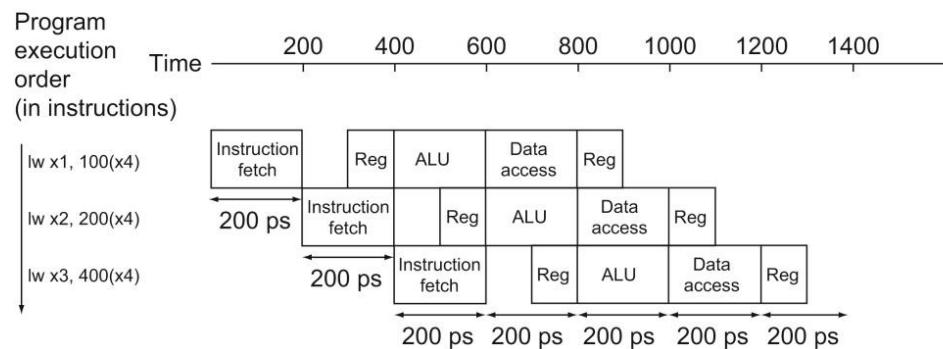
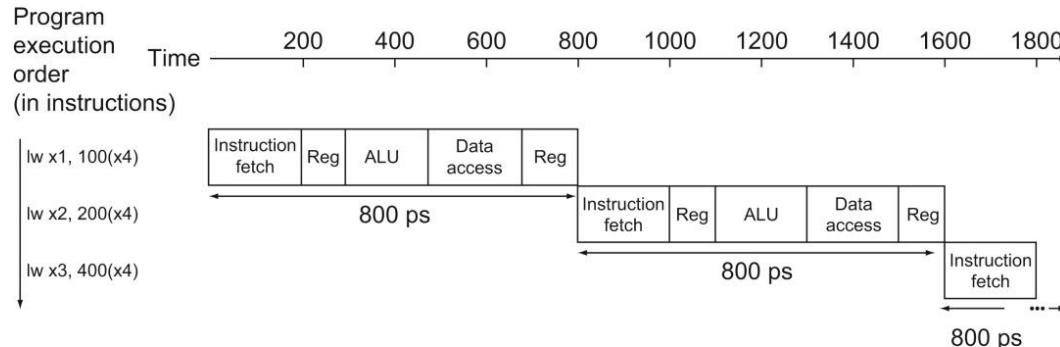


- **Each step takes one clock cycle**
- **One step = one stage**

- **Ifetch: Instruction Fetch**
 - Fetch the instruction from the Instruction Memory
- **Reg/Dec: Registers Fetch and Instruction Decode**
- **Exec: Calculate the memory address**
- **Mem: Read the data from the Data Memory**
- **Wr: Write the data back to the register file**

Pipelining

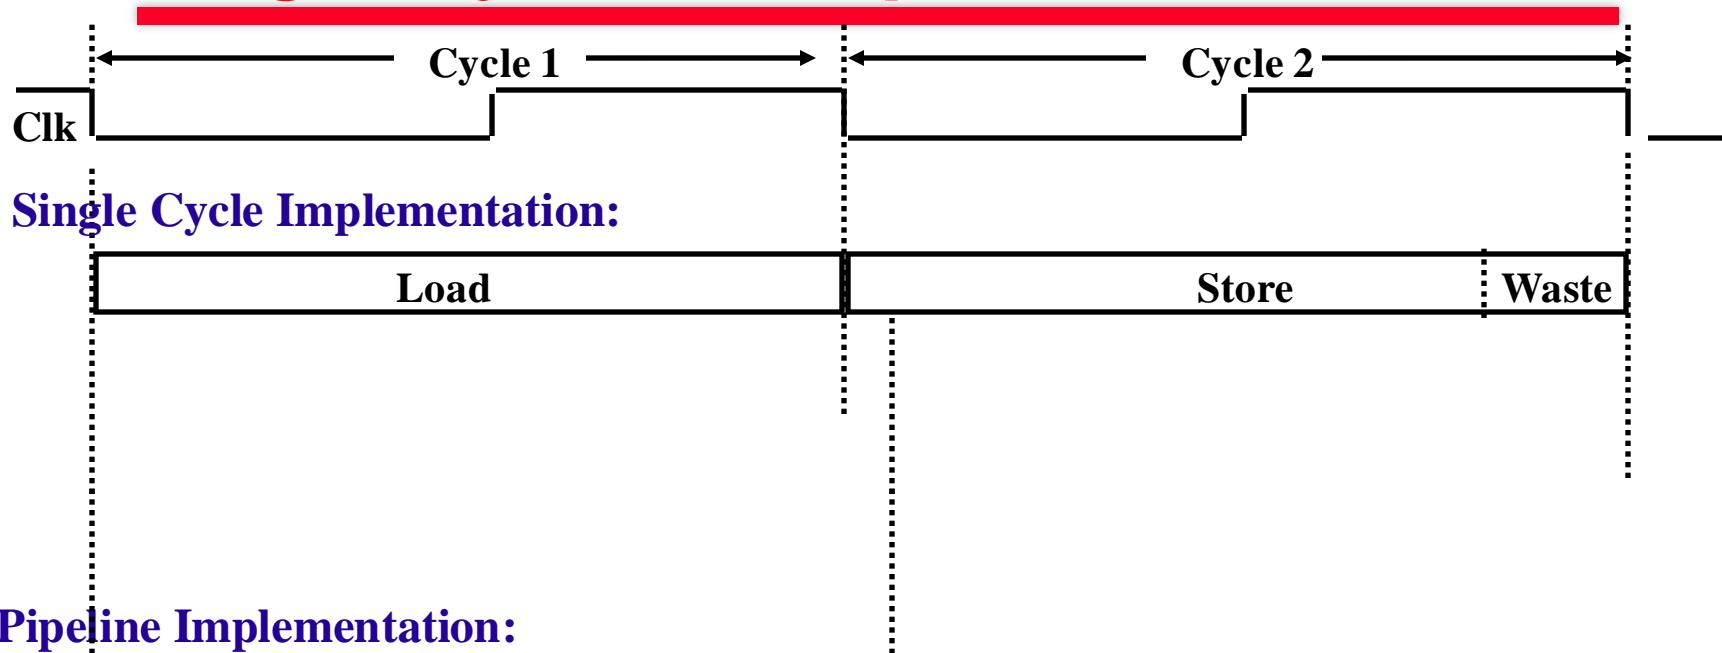
- Improve performance by increasing throughput



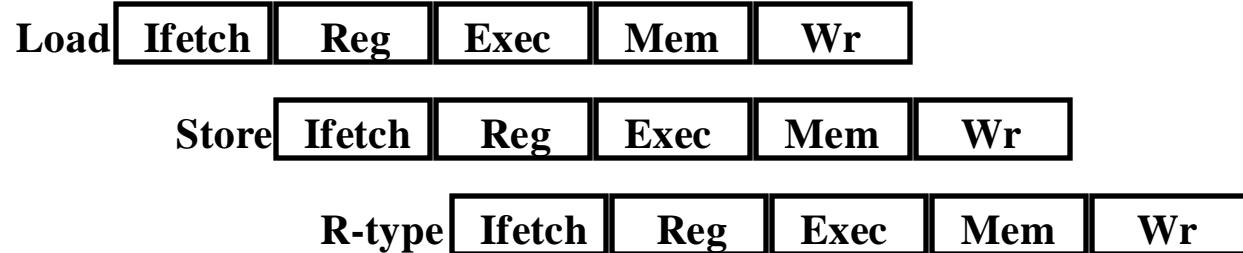
- Do we achieve the ideal speedup (number of stages in the pipeline)? NO!

- pipeline stage time: limited by the slowest resource (ALU operation, or memory access)
- Fill and drain time

Single Cycle vs. Pipeline



Pipeline Implementation:



Why Pipeline? To improve performance!

- **Suppose**

- we execute 100 instructions
 - 45 ns for `Iw` (one cycle)
 - 10 ns for each stage

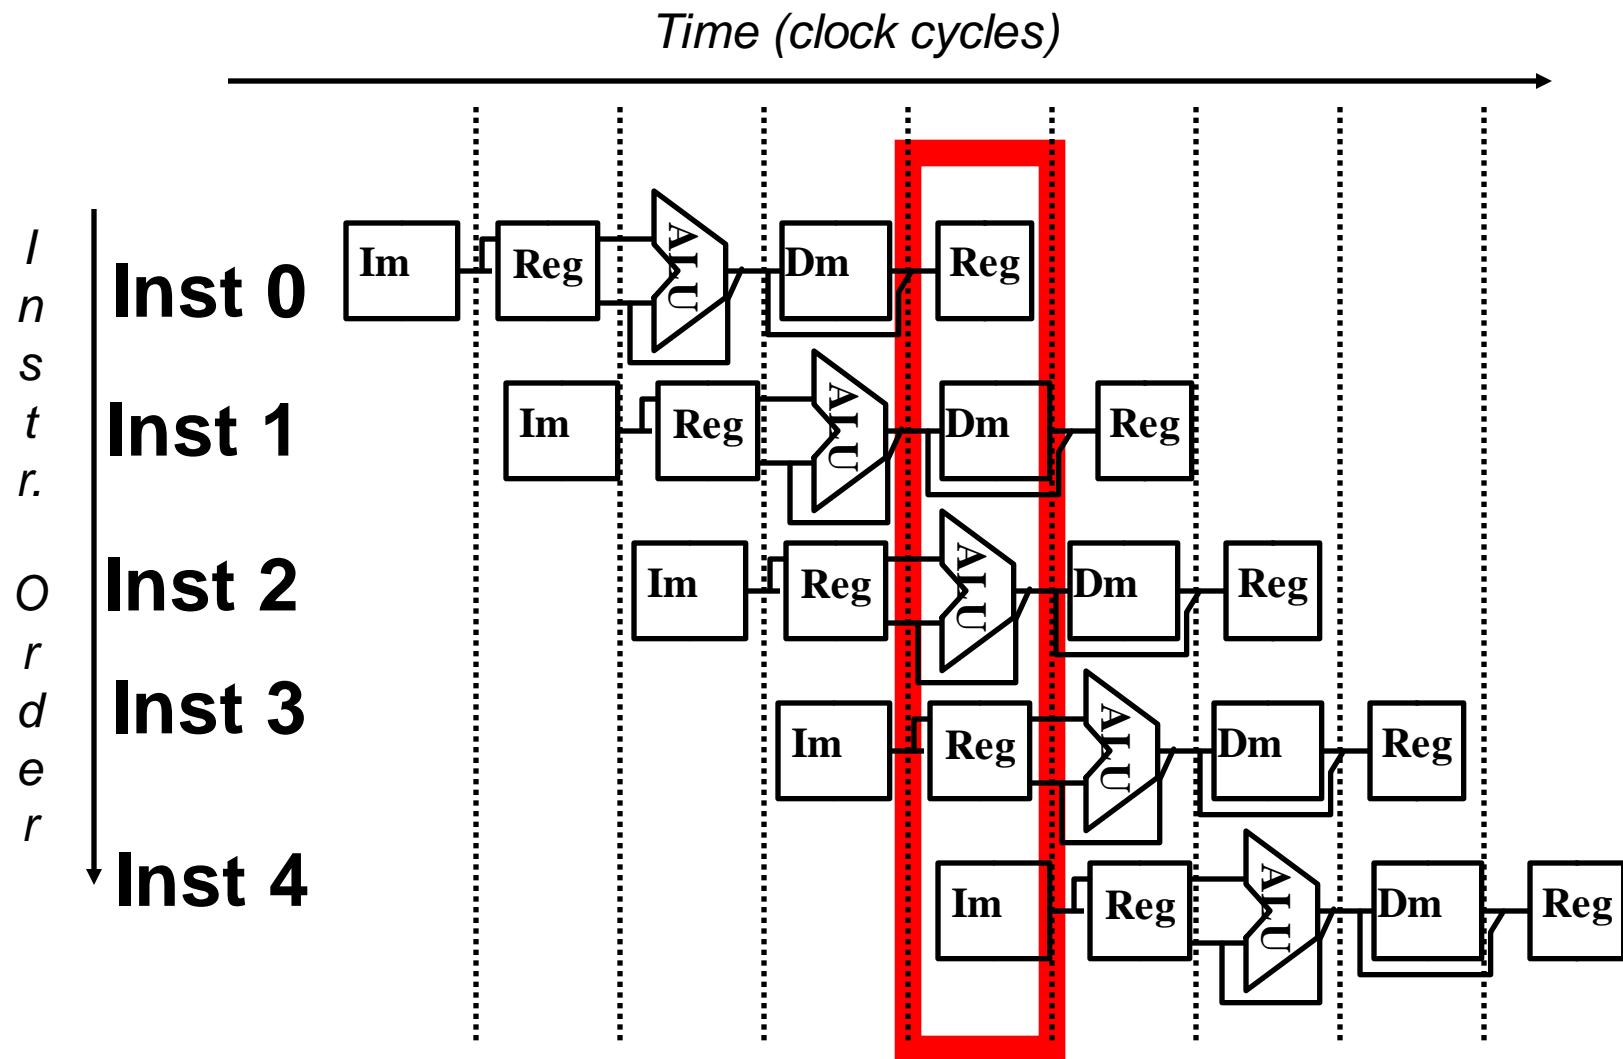
- **Single Cycle Machine**

- $45 \text{ ns/cycle} \times 1 \text{ CPI} \times 100 \text{ inst} = 4500 \text{ ns}$

- **Ideal pipelined machine**

- $10 \text{ ns/cycle} \times (1 \text{ CPI} \times 100 \text{ inst} + 4 \text{ cycle drain}) = 1040 \text{ ns}$

Why Pipeline? Because resources are there!



Can pipelining get us into trouble?

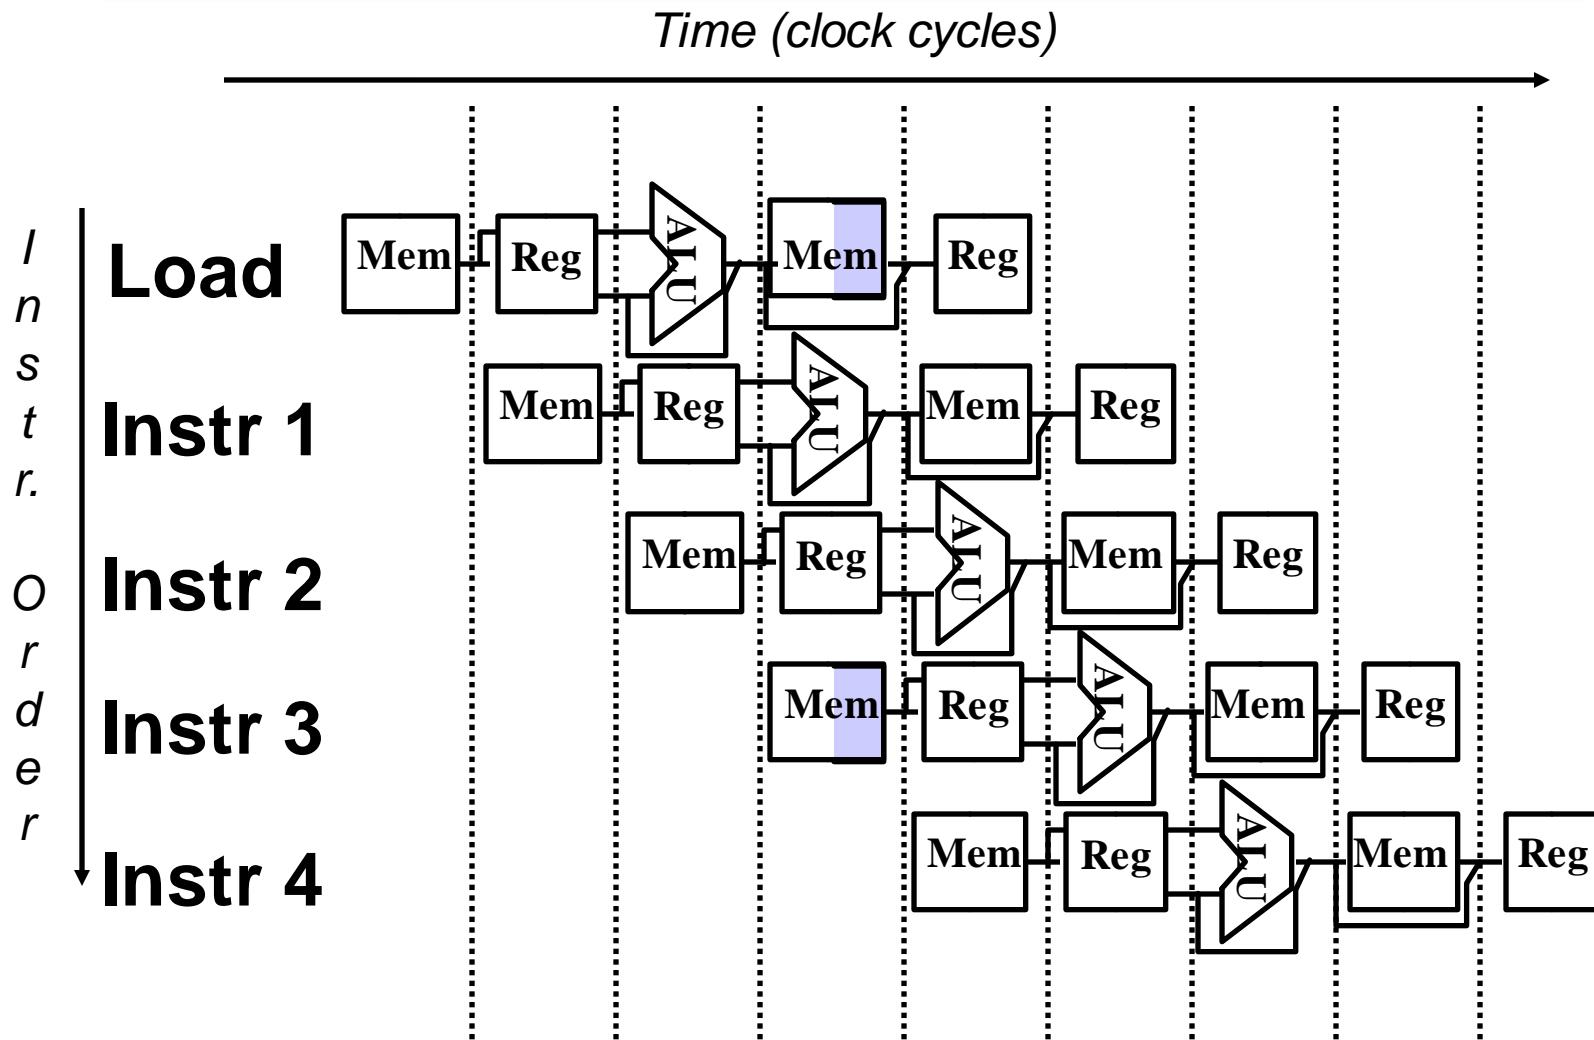
- **Yes: Pipeline Hazards**

- Structural hazards:
 - attempt to use the same resource in two different ways at the same time
 - Control hazards:
 - attempt to make a decision before condition is evaluated
 - Data hazards:
 - attempt to use item before it is ready

- **Can always resolve hazards by waiting**

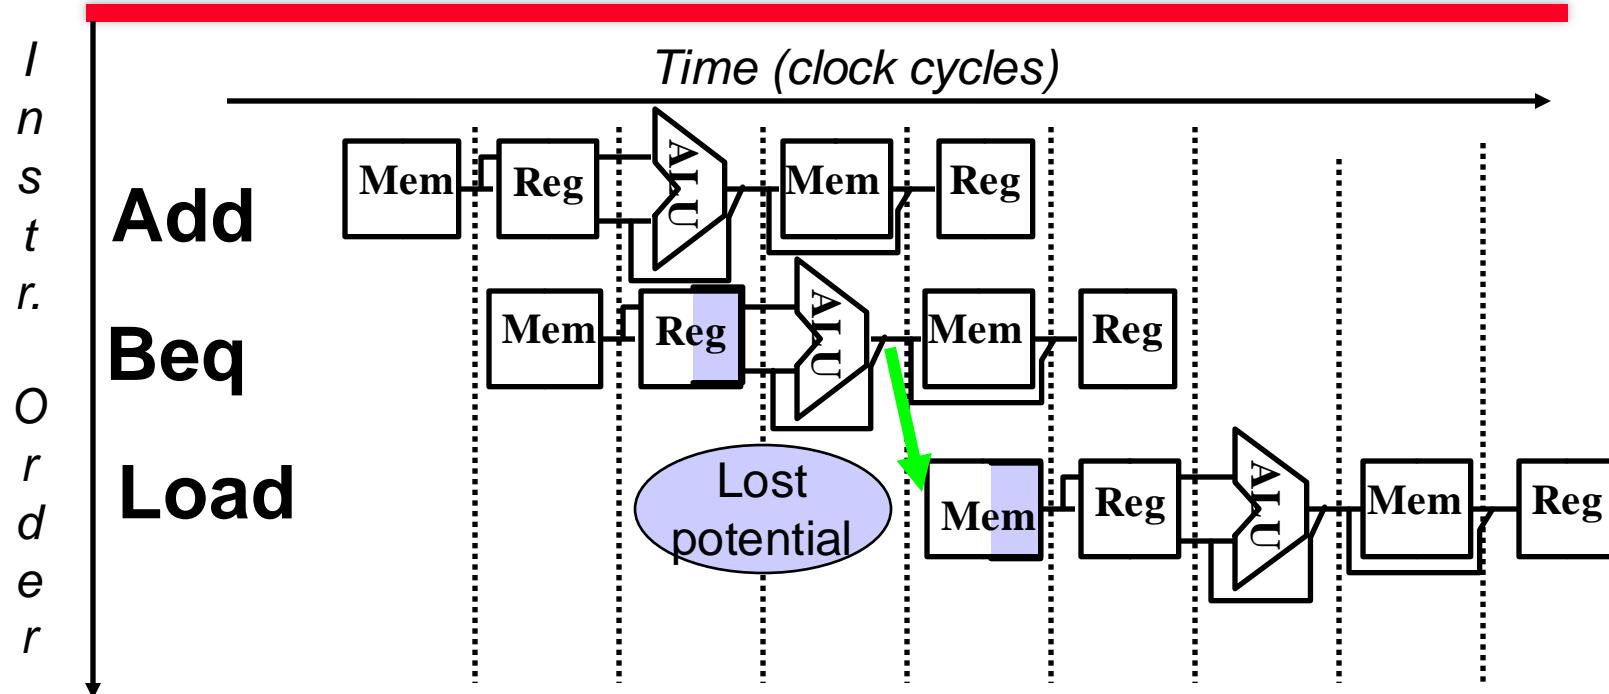
- pipeline control must detect the hazard
 - take action (or delay action) to resolve hazards

Single Memory is a Structural Hazard



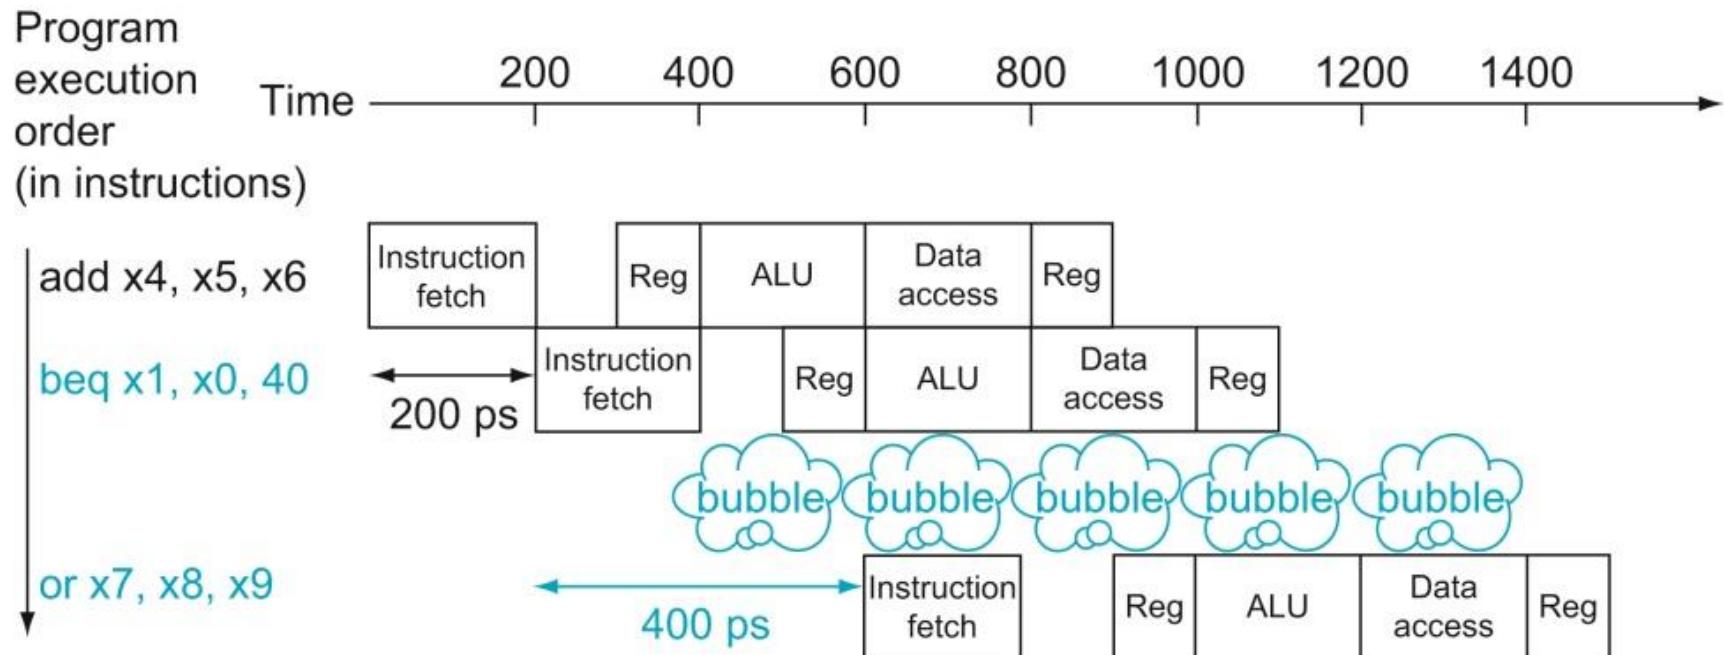
Memory notation: right half highlight means read, left half highlight means write

Control Hazard Solution #1: Stall



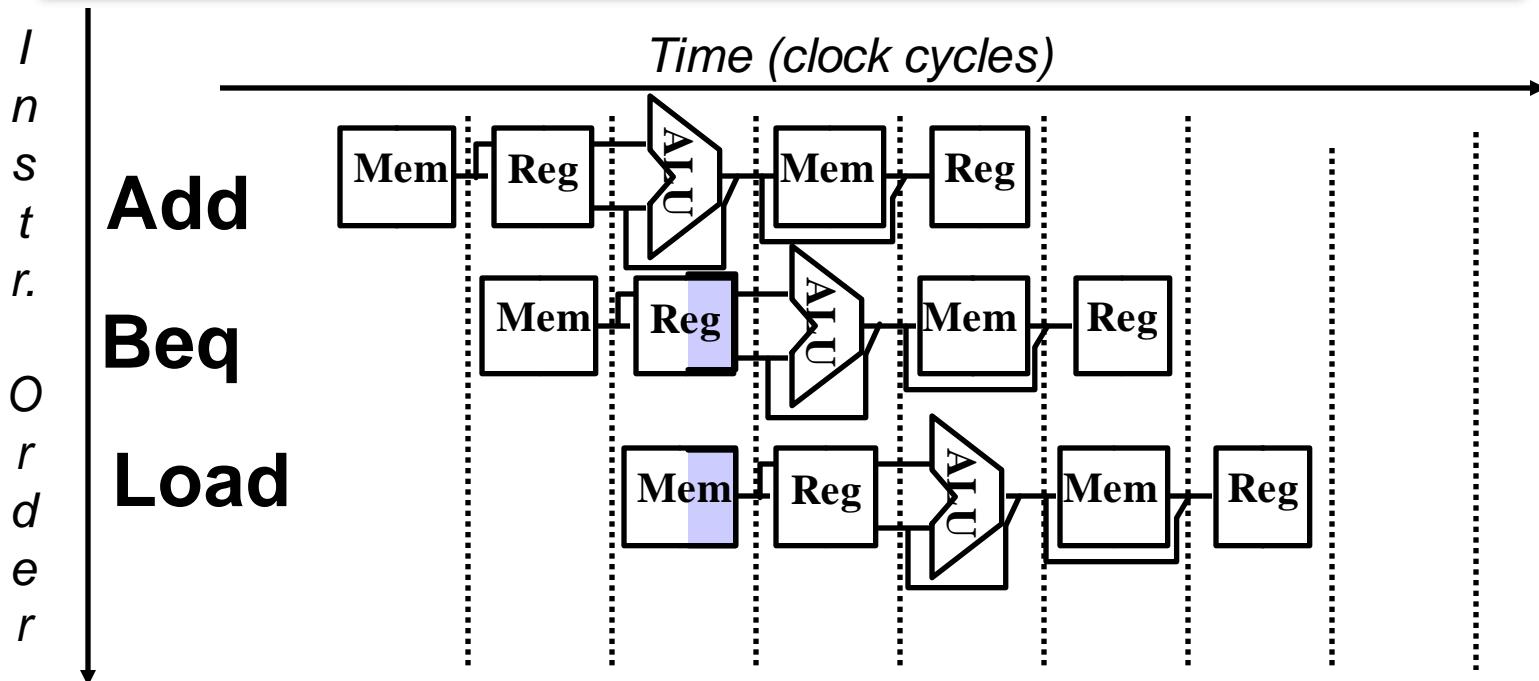
- **Stall: wait until decision is clear**
- **Impact**
 - 2 lost cycles (i.e., 3 clock cycles per branch instruction) => slow
 - Example: If 20% instructions are BEQ, all others have CPI 1, what is the average CPI?

Control Hazard Solution #1: Stall



- Move decision to end of decode by improving hardware
 - save 1 cycle per branch

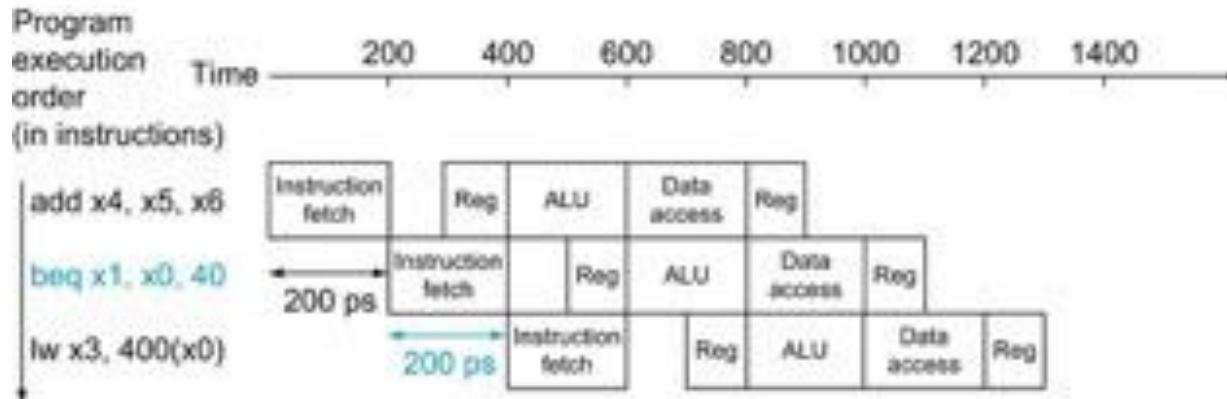
Control Hazard Solution #2: Predict



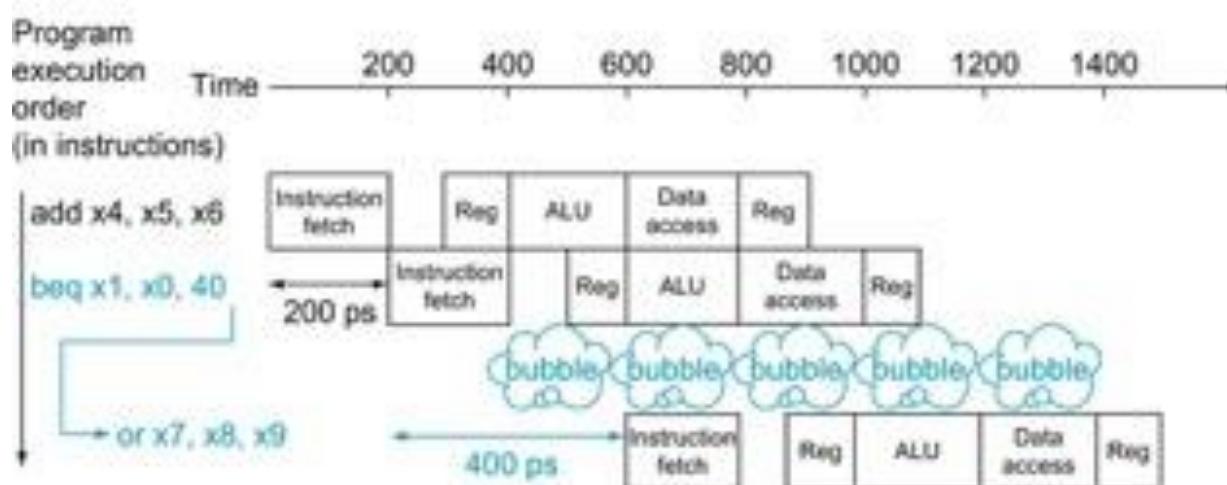
- **Predict:** guess one direction then back up if wrong
- **Impact:** 0 lost cycles per branch instruction if right, 1 if wrong (**prediction accuracy: 50%**)
 - Need to “Squash” and restart following instruction if wrong
 - Produce CPI on branch of $(1 * .5 + 2 * .5) = 1.5$
 - Total CPI might then be: $1.5 * .2 + 1 * .8 = 1.1$ (20% branch)
- **More dynamic scheme:** history of each branch (**prediction accuracy: 90%**)

Control Hazard Solution #2: Predict

Predicting that branches are not taken

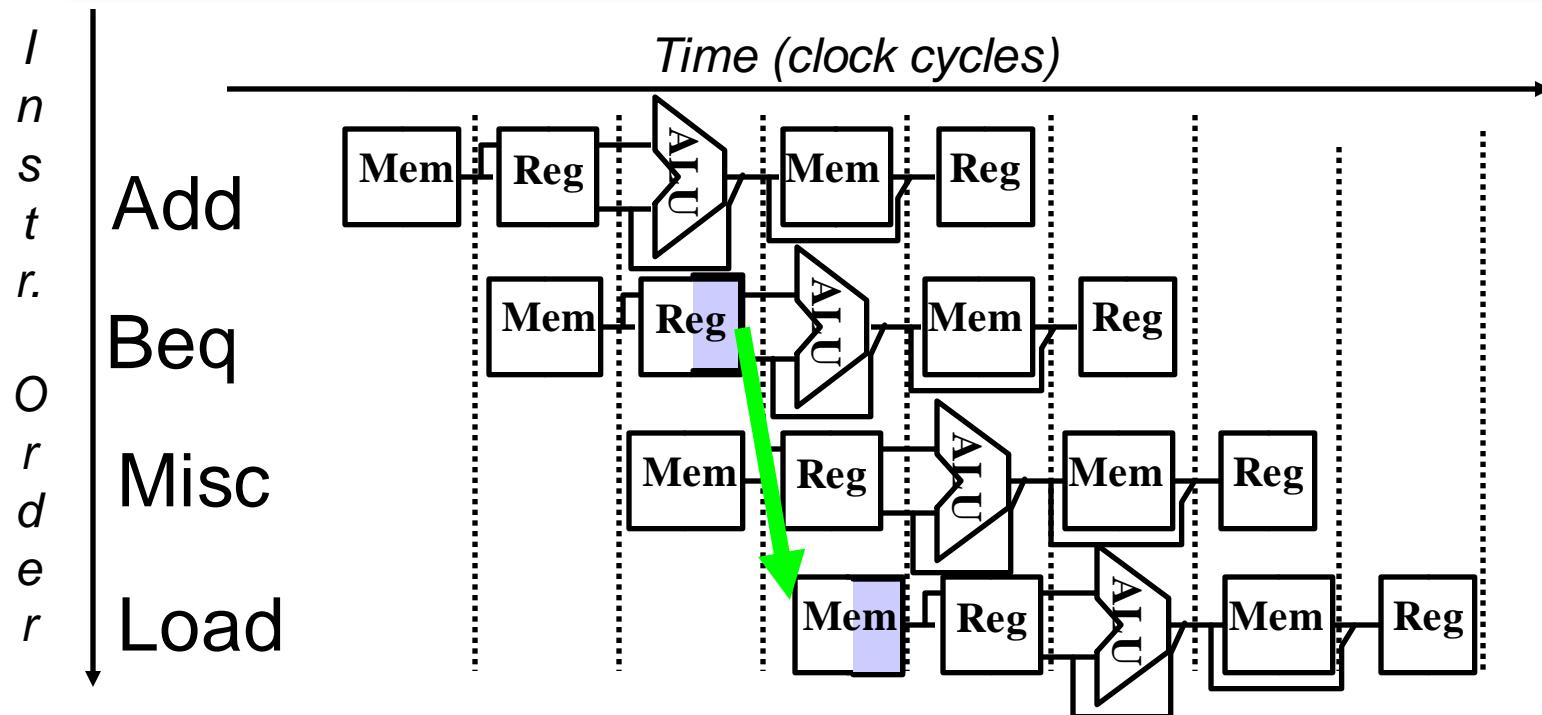


branches not taken



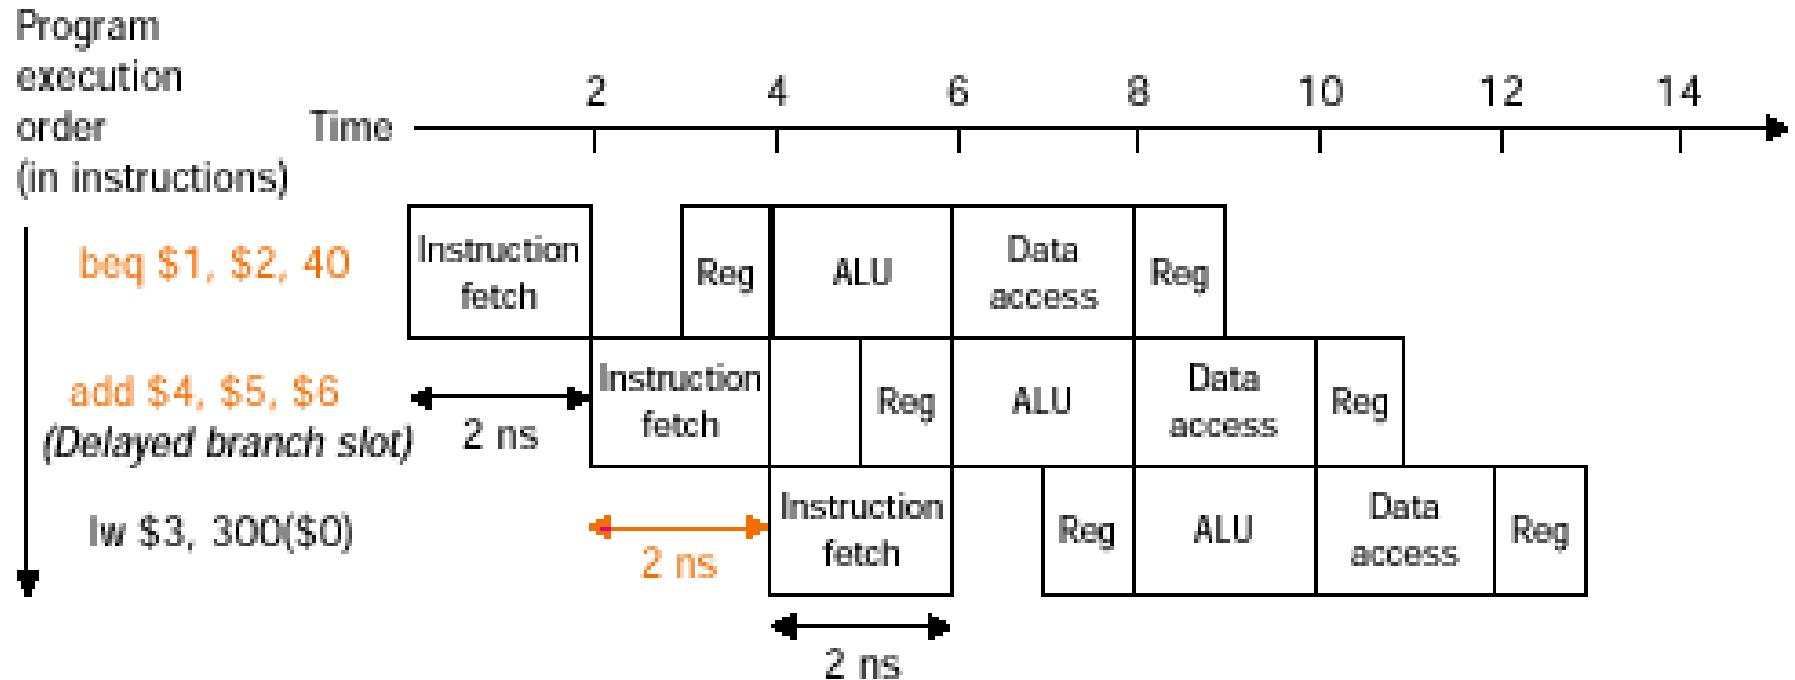
branches taken

Control Hazard Solution #3: Delayed Branch



- **Delayed Branch:** Redefine branch behavior (takes place after next instruction)
- Impact: 0 extra clock cycles per branch instruction if can find instruction to put in “slot” (50% of time)
- The longer the pipeline, the harder to fill
- Used by RISC-V architecture

Control Hazard Solution #3: Delayed Branch



Data Hazard

- An instruction depends on the result of a previous instruction still in the pipeline
- Example:

add r1 ,r2,r3

sub r4, r1 ,r3

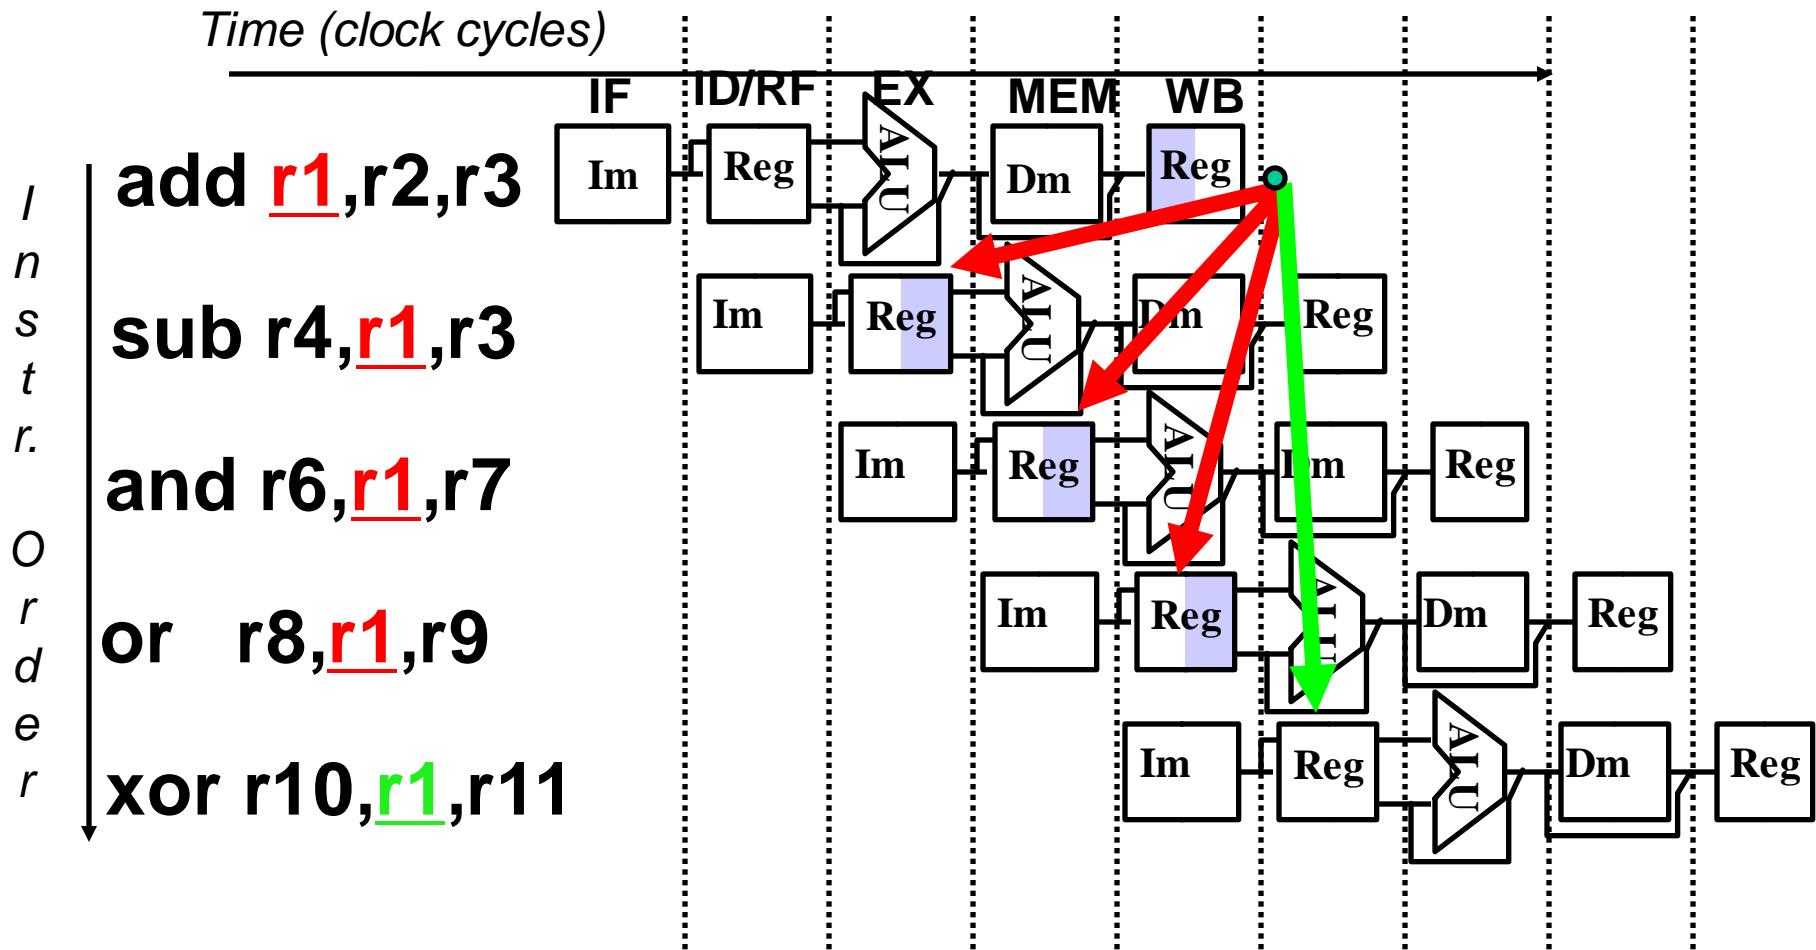
and r6, r1 ,r7

or r8, r1 ,r9

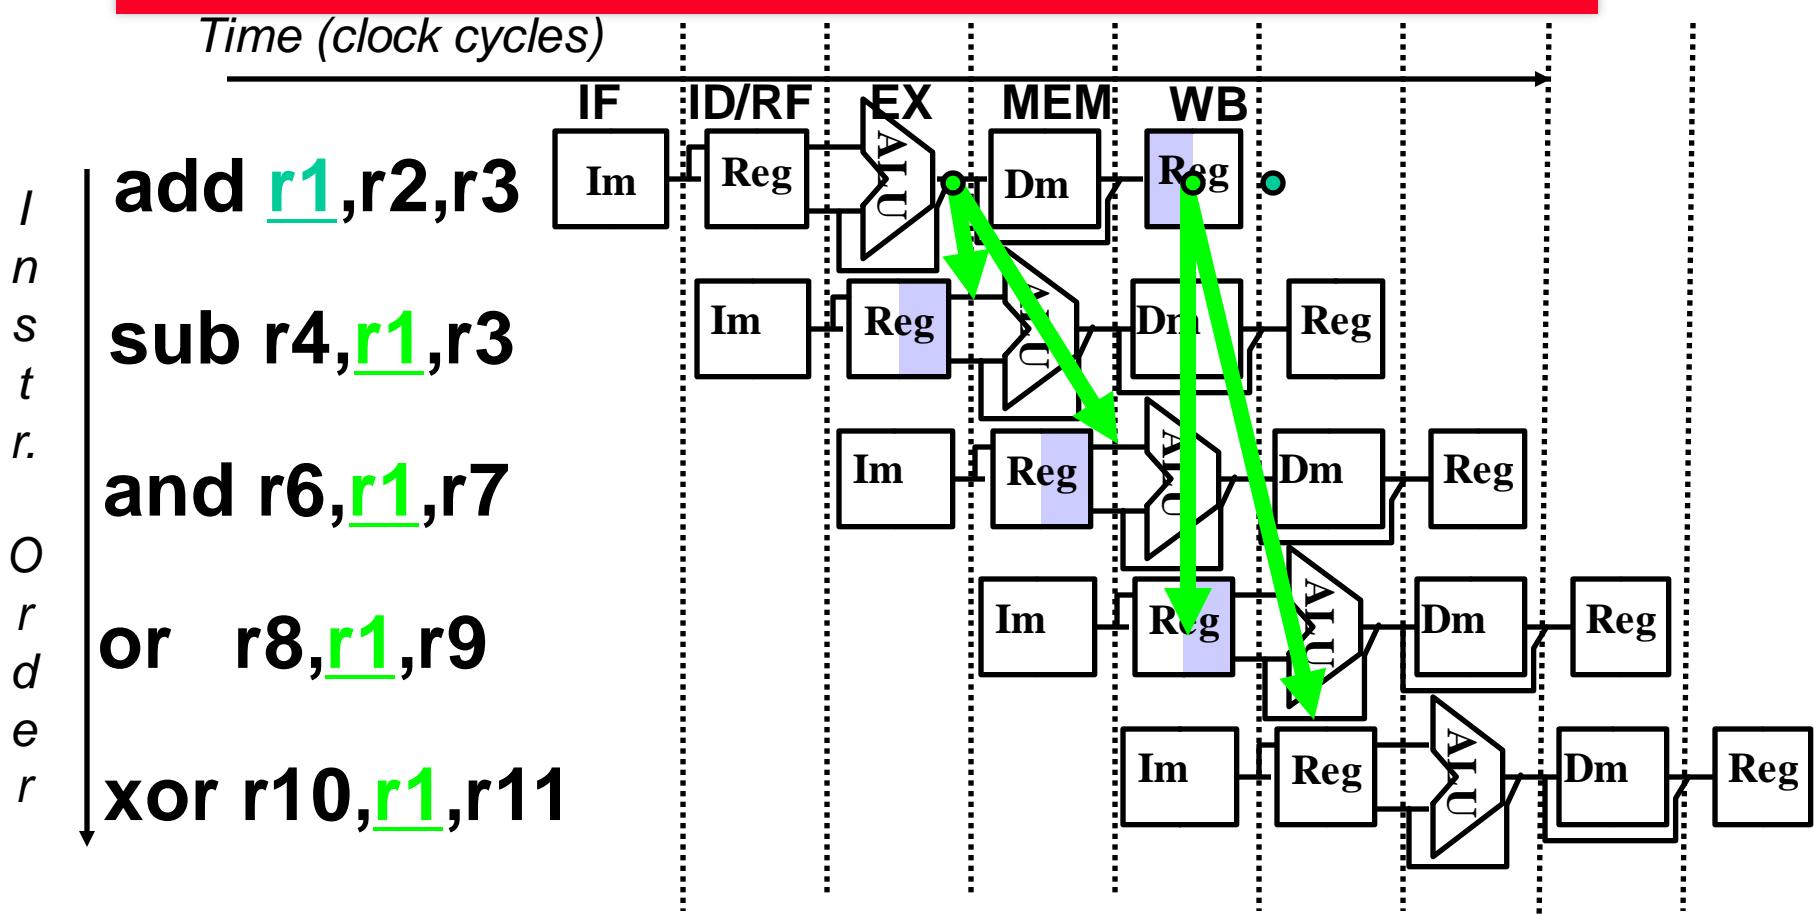
xor r10, r1 ,r11

Data Hazard on r1

- Hazards: Dependencies backwards in time

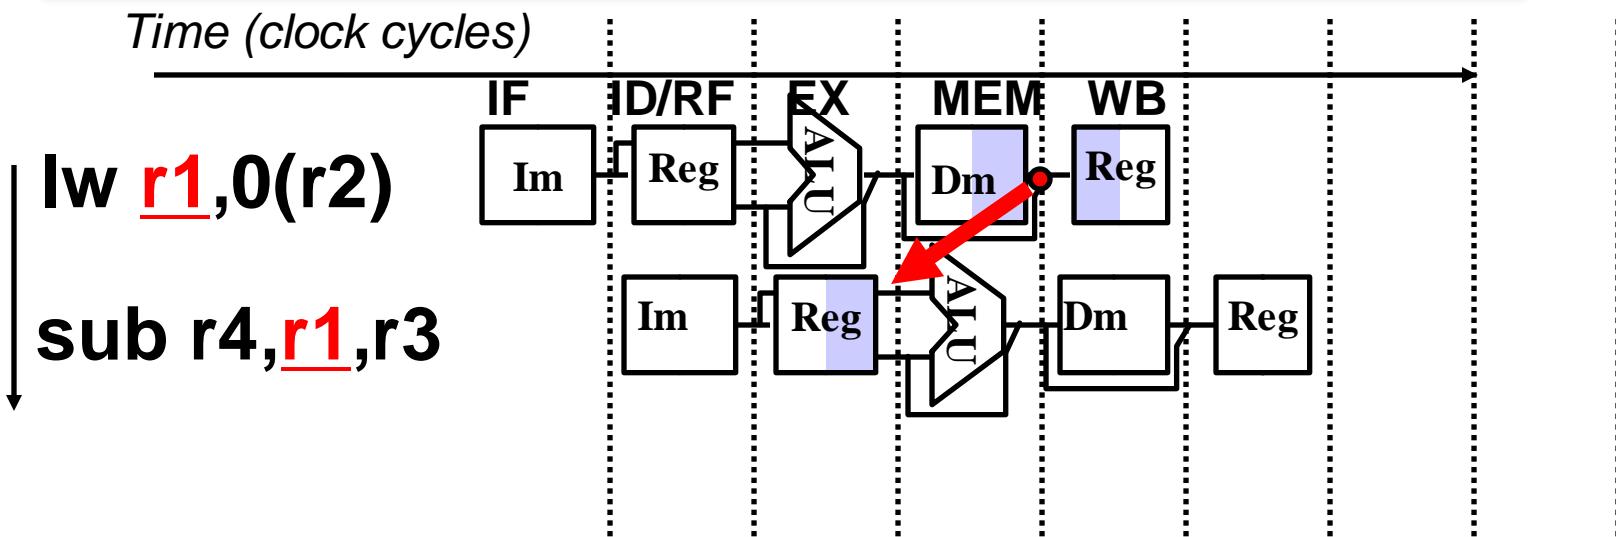


Data Hazard Solution

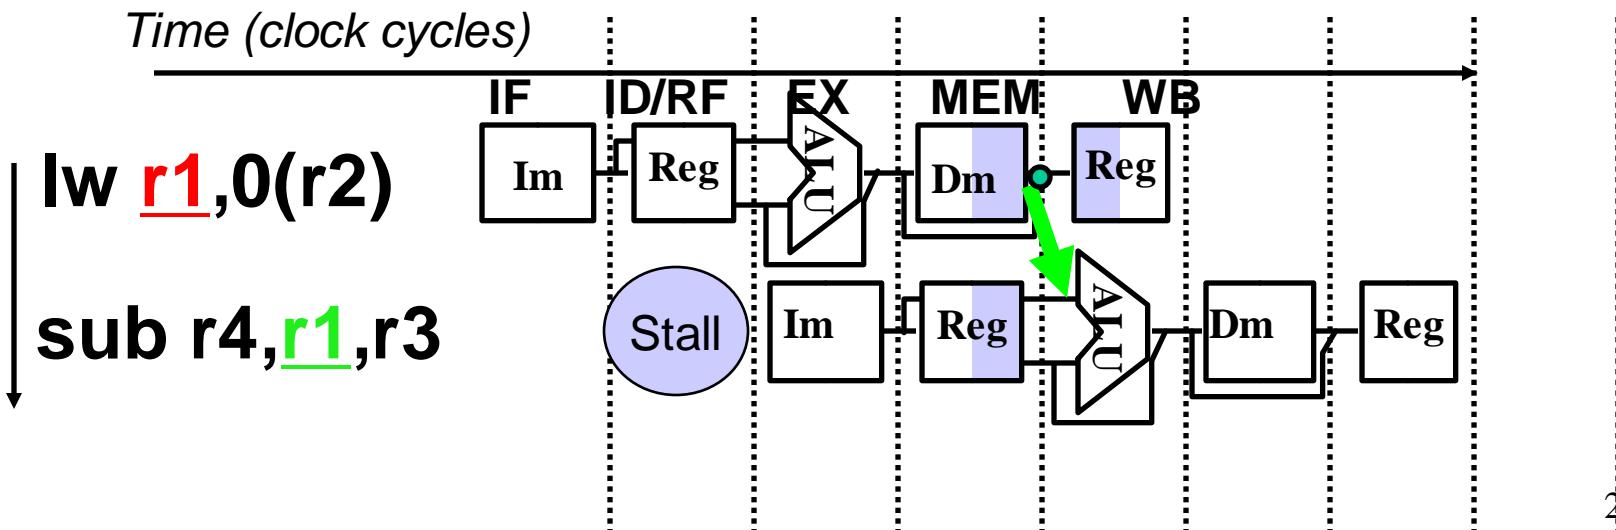


- Forwarding result from one stage to another
 - Adding extra hardware to retrieve the missing item early from the internal resources
- “or” OK if read/write defined properly
- Forwarding can’t prevent all data hazard

What about Loads?

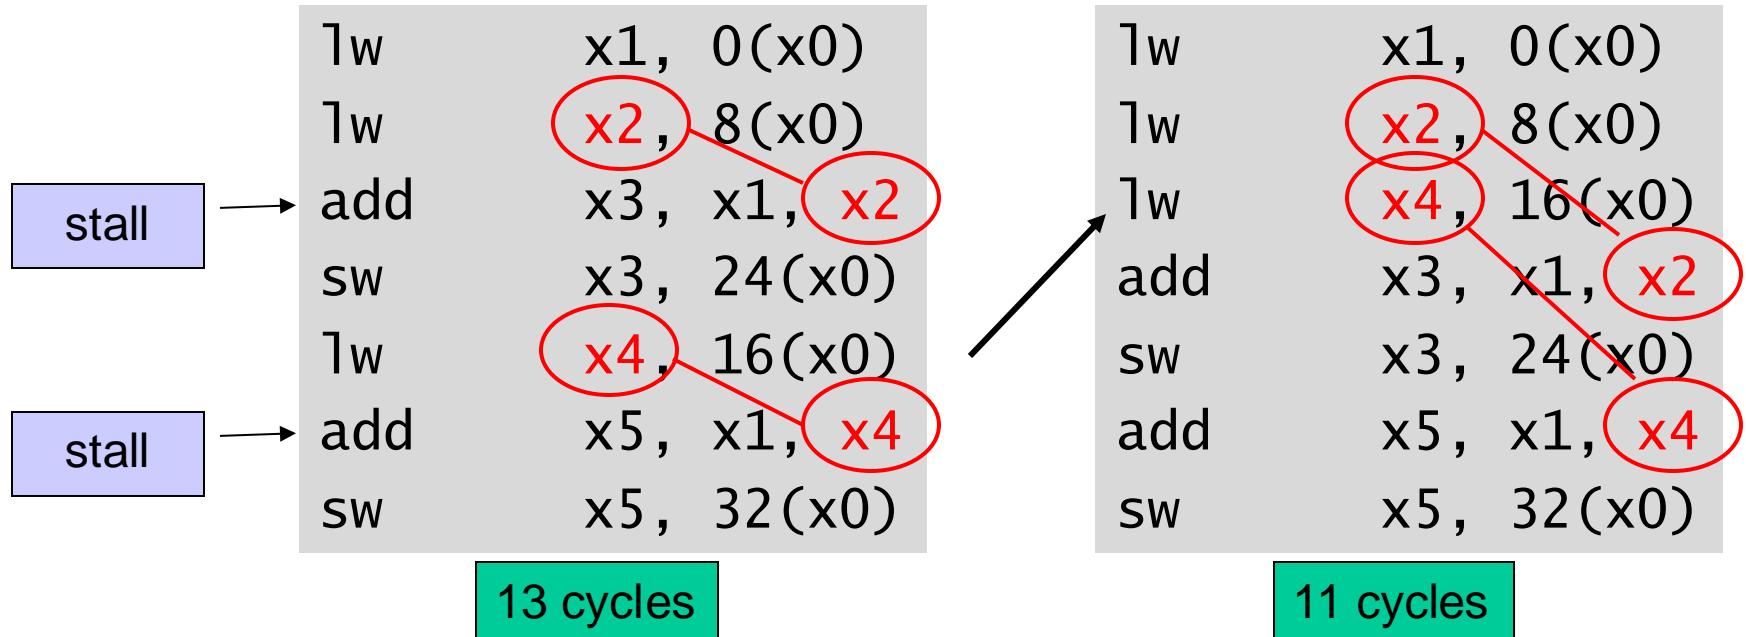


- Can't solve with forwarding
- Must delay/stall instruction dependent on loads



Code Reordering to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for $a = b + e; c = b + f;$



Pipelining Overview Summary

- **Pipelining is a fundamental concept**
 - multiple steps using distinct resources
- **Utilize capabilities of the Datapath by pipelined instruction processing**
 - start next instruction while working on the current one
 - limited by length of longest stage (plus fill/flush)
 - detect and resolve hazards
- **What makes it easy**
 - all instructions are of the same length
 - just a few instruction formats
 - memory operands appear only in loads and stores; Memory addresses are aligned
- **What makes it hard?**
 - structural hazards: suppose we had only one memory
 - control hazards: need to worry about branch instructions
 - data hazards: an instruction depends on a previous one

Pipelining Overview Summary

- In modern processors, what really makes it hard:
 - exception handling
 - trying to improve performance with out-of-order execution, etc.
- We'll build a simple pipeline
 - Only consider the core subset of RISC-V instructions: lw, sw, and, or, add, sub, beq