

# *Topics*


---

- Single Cycle CPU Design (sections 4.1-4.4)
- Pipelining
  - Overview (section 4.6)
  - Pipelined Datapath and control (section 4.7)
  - Data hazard (section 4.8)
  - → Control hazard (section 4.9)


***Thus times do shift,  
each thing his turn does hold;  
New things succeed,  
as former things grow old.***

***Robert Herrick***

# Branch Logic




# Branch Hazards



**branch hazard:** attempt to make a decision before condition is evaluated

# *Branch Hazards*



# *Observations*

---

- **Current design, non-optimized**
  - Branch decision does not occur until MEM stage  $\Rightarrow$  waste 3 CCs
- **Is it possible to reduce branch delay?**
  - YES
  - How? For  $beq\ x, y, \text{label}$ ,  $x \ xor \ y$  then **or all bits**
  - In EXE stage
    - Two CCs branch delay
  - In ID Stage
    - One CC branch delay
- **3 strategies to resolve branch hazard**
  - Delayed branch;
  - Static branch prediction;
  - Dynamic branch Prediction

# *Delayed Branch*

---


- **always execute the instruction following the branch.**
  - Only one will be executed
  
- **Done by compiler or assembler**
  - 50% successful rate

# ***Static Branch Prediction***

---

- Assume the branch will not be taken
- If prediction is wrong, clear the effect of sequential instruction execution.
  - Branch decision is made at MEM stage: instructions in IF, ID, EX stages need to be discarded.
  - Branch decision is made at ID stage: only flush IF/ID pipeline register!
- ***note, flush is not stall! ... we are not repeating the flushed instruction, we are discarding it.***

# *Branch Decision at the MEM Stage*




# *Reducing Penalty*

---

- **Move hardware to an earlier stage: ID stage**
  - Branch Target Address (BTA) adder
  - Register comparator (new!)
  - Control combines Branch and comparator for IF Flush signal, BTA selection in ID stage

# *Branch Decision at the ID Stage*




## *Example: Branch Taken*


---

```
36: sub x10, x4, x8
40: beq x1, x3, 16 # 16 is the value in the immediate
44: and x12, x2, x5
48: or x13, x2, x6
52: add x14, x4, x2
56: slt x15, x6, x7
...
72: lw x4, 48(x7)
```

# Example: Branch Taken




# Example: Branch Taken



# *Dynamic Branch Prediction*


- **Static branch prediction is crude!**
- **Take history into consideration**
  - If a branch was taken last time, then fetch the new instruction from the same place
  - Branch prediction buffer
    - contains a bit (or bits) which tells whether the branch was recently taken or not
  - Is the prediction correct? Any bad effect?
  - 1-bit prediction scheme
  - 2-bit prediction scheme



# Exercise

---

- Consider a loop branch that branches nine times in a row, then is not taken once. What is the prediction accuracy for this branch, assuming the prediction bit for this branch remains in the prediction buffer?
  - 1-bit prediction?
  - 2-bit prediction?
- Actual: T, T ..., T, NT
- 1-bit prediction:  
NT, T...T, T (80% accuracy)
- 2-bit prediction?  
NT, NT, T...T, T (70% accuracy)



# *Exceptions*

---

- **Exceptions: events other than branch or jump that change the normal flow of instruction**
  - Arithmetic overflow, undefined instruction, etc
  - Internal of the processor
  - Interrupts from external – I/O interrupts
- **Use arithmetic overflow as an example**
  - When an overflow is detected, we need to transfer control to the exception handling routine immediately
    - RISC-V uses location `0x 1C09 0000`
    - Detected in the EX stage
  - Similar idea as branch hazard
    - De-assert all control signals in EX and ID stages, flush IF/ID

# *Example*


---

```
sub  x11, x2, x4
and  x12, x2, x5
or   x13, x2, x6
add  x1, x2, x1          -- overflow occurs
sub  x15, x6, x7
lw   x16, 100(x7)
```


## Exceptions handling routine:

0x 1C09 0000	sw x25, 1000(x0)
0x 1C09 0004	sw x26, 1004(x0)

# Clock 6



# Clock 7



Clock 7

# *Performance of Pipelined Datapath*

---

- **Functional unit time**
  - 200ps for memory access
  - 100ps for ALU operation
  - 50ps for register file access
- **Instruction frequencies**
  - 25% loads,
  - 10% stores,
  - 45% R-type
  - 20% branches,
- **For pipelined datapath, compute average CPI and CCT**
  - 50% of load are immediately followed by an instruction that uses the result
  - Branch delay on misprediction is 1 clock cycle and 25% branches are mispredicted
  - Assume branch condition evaluation is done in the ID stage

# **Solution**

---

- **For pipelined design**

- CPI for lw:  $1 \times 0.5 + 2 \times 0.5 = 1.5$
- CPI for branch:  $1 \times 0.75 + 2 \times 0.25 = 1.25$
- Average CPI =  $1.5 \times 25\% + 1 \times 10\% + 1 \times 45\% + 1.25 \times 20\% = 1.175$
- CCT = 200ps

# *Example 1*

---

Consider the following code:

```
addi x10, x0, 2  
loop: addi x10, x10, -1  
       addi x9, x10, -1  
       bne x10, x0, loop  
       add x11, x10, x9
```

For all parts, assume forwarding

- How many cycles does this code take if there is no branch prediction, and branches are resolved (i.e., we know if they are taken or not taken) at the end of the ID stage?
- How many cycles does this code take if branches are predicted not taken, and branches are resolved at the end of the ID stage?
- How many cycles does this code take if branches are predicted not taken, and branches are resolved at the end of the MEM stage?

# Example 1a Answer

---

a) 14

Ins	CC1	CC2	CC3	CC4	CC5	CC6	CC7	CC8	CC9	CC10	CC11	CC12	CC13	CC14
I1	IF	ID	EX	M	WB									
I2		IF	ID	EX	M	WB								
I3			IF	ID	EX	M	WB							
I4				IF	ID	EX	M							
I2						IF	ID	EX	M	WB				
I3							IF	ID	EX	M	WB			
I4							IF	ID	EX	M				
I5										IF	ID	EX	M	WB

## *Example 1b Answer*

---

b) 13

The second branch's stall will now be avoided as I'm predicting not taken and predicting correct. So, execution time = 14 - 1 = 13 cycles.

# Example 1c Answer

---

c) 15

The second branch will not cause any stalls as it is a

Ins	CC1	CC2	CC3	CC4	CC5	CC6	CC7	CC8	CC9	10	11	12	13	14	15
I1	IF	ID	EX	M	WB										
I2		IF	ID	EX	M	WB									
I3			IF	ID	EX	M	WB								
I4				IF	ID	EX	M								
I2								IF	ID	EX	M	WB			
I3									IF	ID	EX	M	WB		
I4									IF	ID	EX	M			
I5										IF	ID	EX	M	WB	

## Example 2

---

- Schedule the following code for the pipelined processor with forwarding and reduced branch delay using a single branch delay slot to minimize stall cycles:

```
loop: lw x1,0(x2)      # x1 array element
      add x1, x1, x3    # add constant in x3
      sw x1,0(x2)      # store result array element
      addi x2, x2, -4   # decrement address by 4
      bne x2, x4, loop  # branch if x2 != x4
```

- Assuming the initial value of  $x2 = x4 + 40$  (i.e., it loops 10 times); assuming condition is evaluated at ID stage
- What is the CPI and total number of cycles needed to run the code with and without compiler scheduling?

# Example 2 Solution

- Without compiler scheduling

loop: lw x1,0(x2)

stall

add x1, x1, x3

sw x1,0(x2)

addi x2, x2, -4

stall

bne x2, x4, loop

stall

Ignoring the initial 4 cycles to fill the pipeline:

Each iteration takes = 8 cycles

CPI = 8/5 = 1.6

Total cycles = 8 x 10 = 80 cycles

- With compiler scheduling

loop: lw x1,0(x2)

addi x2, x2, -4

add x1, x1, x3

bne x2, x4, loop

**sw x1, 4(x2)**

Move between  
lw add

Move to  
branch delay slot

Adjust  
address  
offset

Ignoring the initial 4 cycles to fill the pipeline:

Each iteration takes = 5 cycles

CPI = 5/5 = 1

Total cycles = 5 x 10 = 50 cycles

Speedup = 80/ 50 = 1.6