

Topics for Chapter 2

- **Part 1**

- RISC-V Instruction Set (Chapter 2.1)
- Arithmetic instructions (Chapter 2.2)
- Introduction to RARS
- Memory access instructions (Chapters 2.3)
- Bitwise instructions (Chapter 2.6)
- Decision making instructions (Chapter 2.7)
- Arrays vs. pointers (Chapter 2.14)

- **Part 2:**



- → Procedure Calls (Chapters 2.8)

- **Part 3:**

- RISC-V Instruction Format (Chapter 2.5)
- RISC-V Addressing Modes (Chapter 2.10)

Stored Program Concept

The BIG Picture



- Instructions and data are stored in memory.
- Each register stores how many bits? how many bytes? how many words?
- How much data can I access from memory at a time?
- Each instruction is how many bits?
- Memory addresses are how many bits?
- RISC-V memory capacity in bytes? in words?


RISC-V Memory Allocation

SP → 0000 003f ffff fff0_{hex}

In RARS, sp initial value=0x7fff effc

0000 0000 1000 0000_{hex}

PC → 0000 0000 0040 0000_{hex}



- **Static:** Variables declared once per program, cease to exist after execution completes. e.g., C globals, arrays
- **Heap:** Variables declared dynamically (malloc)
- **Stack:** Space to be used by procedure during execution; this is where we can save register values

RISC-V Memory Allocation

Text segment

- Program code
- Addresses 0x0040 0000 to 0xFFFF FFFF

Data segment


- Addresses 0x1000 0000 to 0xFFFF FFFF

SP → 0000 003f ffff fff0_{hex}


0000 0000 1000 0000_{hex}

PC → 0000 0000 0040 0000_{hex}

0



Processor



The Program Counter (PC)

- **Special purpose register**
- **Not available to programmer**
- **Holds address of?**
- **Assuming currently executing the first instruction**
 - $(5048)_{10} = 0x13B8$
 - PC = 0x0000 _____
 - PC + 4 = 0x0000 _____
 - Byte stored at 0x0000 13BA?

| Memory Location | Instruction |
|--------------------|-------------------------|
| 5048 | 0x00000413 |
| 5052 | addi s1, x0, 1 |
| 5056 | beq s1, s0, continue |
| 5060 | jal x0, exit |
| 5064 | |
| 5068 (continue) | |
| ... | |
| 6000 (exit) | |

This is the machine representation of:
addi s0, x0, 0

Procedure Call in C Program

```
main() {  
    int i,j,k,m;  
    ...  
    i = mult(j,k); ...  
    m = mult(i,i); ...  
}  
  
/* really dumb mult function  
 */  
  
int mult (int mcand, int  
mlier){  
    int product;  
    product = 0;  
    while (mlier > 0) {  
        product = product + mcand;  
        mlier = mlier -1; }  
    return product;  
}
```

What information must a compiler/programmer keep track of?

What instructions can accomplish this?

Procedure Call Bookkeeping

- **Use registers**
- **Register conventions:**
 - Return address ra
 - Arguments a0, a1, a2, ..., a7
 - Return value a0, a1
 - Local variables s0, s1, ..., s11
 - Temporary variables t0, t1, ..., t6
- **More variables?**
 - Spill registers: use the stack in memory

RISC-V Support for Function Calls

```
... sum(a,b); ... /* a,b:s0,s1 */          C
}
int sum(int x, int y) {
    return x+y;
}
```

| address | RISC-V |
|---------|-----------------------------------|
| 1000 | add a0,s0,zero # x = a |
| 1004 | add a1,s1,zero # y = b |
| 1008 | jal ra, sum # new instruction |
| 1012 | ... |
| 2000 | sum: add a0,a0,a1 |
| 2004 | jalr zero, ra,0 # new instruction |

RISC-V Instruction: jal

- Syntax for `jal` (jump and link)

- `jal ra, label`

- **`jal` should really be called `la j` for “link and jump”:**

- Step 1 (link): Save address of *next* instruction (i.e., $PC+4$) into `ra`
 - Why next instruction? Why not current one?
 - Step 2 (jump): Jump to the given label

- Why `jal`?

- `ra` automatically saves $PC+4$
 - No need to know where the code is loaded into memory
 - Make the common case (function calls) fast

RISC-V Instruction: jalr

- **Syntax for jalr (jump register):**

jalr saved_addr, jump_addr, imm

- saved_addr: where curr instruction address+4 will be stored
- jump_addr: contains address to jump to
- imm: added to jump_addr (always 0 for our needs)

- **Why use jalr?**

- a function might be called many times, so we can't return to a fixed place.

Use of jalr

- **ret and jr pseudo-instructions**

jalr x0, ra, 0

- **Call function at any 32-bit absolute address**

lui x1, <hi20bits>

jalr ra, x1, <lo12bits>

- **Jump PC-relative with 32-bit offset**

auipc x1, <hi20bits>

jalr x0, x1, <lo12bits>

RISC-V Instructions for Function Calls

| | | | | | |
|------|-------------------|---|---------|-----|---------------------------|
| jal | Jump And Link | J | 1101111 | | rd = PC+4; PC += imm |
| jalr | Jump And Link Reg | I | 1100111 | 0x0 | rd = PC+4; PC = rs1 + imm |

Steps for Making a Procedure Call (V. 1)

- 1) Place parameter in registers → `a0-a7`
- 2) Call procedure → `jal ra, PROC`
ra is x1
- 3) Perform procedure's operations
- 4) Place result in register for caller → `a0, a1`
- 5) return to caller → `jalr x0, ra, 0`

Caller

jal ra, ProcedureLabel

| Address | Instruction |
|---------|------------------------|
| 4060 | add ... |
| 4064 | ori ... |
| 4068 | and ... |
| 4072 | jal ra, ProcedureLabel |
| 4076 | sll ... |

→ ra = 4076

| Address | Instruction |
|---------|-------------|
| 5044 | add... |

ProcedureLabel:

Callee

jalr zero, ra, 0

| Address | Instruction |
|---------|------------------|
| 5060 | srl... |
| 5064 | jalr zero, ra, 0 |

→ 4076

| Address | Instruction |
|---------|-------------|
| 4076 | sll ... |

- jal ra, LABEL stores PC+4 in ra
 - Used by caller
- jalr x0, ra, 0 uses ra to modify PC
 - Used by callee

Procedure Call Example: C Code

```
main() {  
    int i,j,k,m; /* i-m:s0-s3 */  
    ...  
    i = mult(j,k); ...  
    m = mult(i,i); ...  
}  
  
int mult (int mcand, int mlier){  
    int product;  
  
    product = 0;  
    while (mlier > 0) {  
        product += mcand;  
        mlier -= 1; }  
    return product;  
}
```

What are the arguments?

What are the results?

Procedure Call Example: main function

main:

```
addi a0,s1,0          # arg0 = j
addi a1,s2,0          # arg1 = k
jal ra, mult          # call mult
addi s0,a0,0          # i = mult()

...
# arg0 = i, already in a0
addi a1,s0,0          # arg1 = i
jal ra, mult          # call mult
addi s3,a0,0          # m = mult()

...
addi a7, zero,10
ecall
```

```
main() {
    int i,j,k,m; /* i-m:s0-s3 */
    ...
    i = mult(j,k); ...
    m = mult(i,i); ...
}
```

Procedure Call Example: sub function

mult:

```
addi t0,zero,0          # prod=0
```

Loop:


```
bge 0,a1,Fin          # if mlr <= 0, goto Fin
add t0,t0,a0            # prod+=mc
addi a1,a1,-1           # mlr-=1
jal zero,Loop           # goto Loop
```

Fin:

```
addi a0,t0,0            # a0=prod
jalr zero, ra, 0        # return
```

```
int mult (int mcand, int mlier) {
    int product = 0;
    while (mlier > 0)  {
        product += mcand;
        mlier -= 1; }
    return product;
}
```

What Do We Know So Far?



- › In RISC-V, "argument" registers for passing parameters are?
- › jal instruction does two actions, those are?
- › RISC-V procedure returns values in?
- › jalr zero, ra, 0 jumps to address in ra, which is?


Nested Procedures

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```

- Something called `sumSquare`, now `sumSquare` is calling `mult`.
- `ra` may be overwritten
 - Need to save `sumSquare` return address before call to `mult`.
- In general, may need to save some other info in addition to `ra`.

Using the Stack to Store Other Info

- register **sp** : the last used space in the stack.
- When **a0-a7** are not enough
 - Spill arguments to the stack
- If the caller uses **s0** before and after the procedure call and the callee also uses **s0**
 - Preserve registers in stack
- **Stack grows downwards**



Stack Pointer (sp)

- Available to programmer
- Allocate space on stack
 - Decrement sp
 - Grow from top down
- Use space on stack
 - Push: sw reg, offset (sp)
 - Pop: lw reg, offset (sp)
- Release space on stack
 - Increment sp
 - Shrink towards top

```
square_sum:  
    addi sp, sp, -4  
    sw ra, 0(sp)
```

...

```
    lw ra, 0(sp)  
    addi sp, sp, 4  
    jalr zero, ra, 0
```

Popping does not make the data disappear.
Neither does releasing the space!

Steps for Making a Procedure Call (V. 2)

- 1) Caller saves **necessary** values onto stack.
- 2) Assign argument(s) in a0-a7, if any.
- 3) `jal ra, proc` -- call procedure
- 4) Perform procedure's operations
- 5) Place result in a0, a1 if any, for caller
- 6) `jalr x0, ra, 0` – return to caller
- 7) Caller restore values from stack.

Using the Stack: Example

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y; }
```

sumSquare:

“push” addi sp,sp,-8 # space on stack
 sw ra, 4(sp) # save ret addr
 sw a1, 0(sp) # save y

 add a1,a0,zero # mult(x,x)
 jal ra, mult # call mult

“pop” lw a1, 0(sp) # restore y
 add a0,a0,a1 # mult() + y
 lw ra, 4(sp) # get ret addr
 addi sp,sp,8 # restore stack
 jalr zero, ra, 0

mult: ...

Rules for Procedures

- Called with a `jal` instruction, returns with a `jalr zero,ra,0`
- Accepts up to 7 arguments in `a0, a1, a2, ..., a7`
- Return value is always in `a0` (and if necessary in `a1`)
- Must follow **register conventions** (even in functions that only you will call)!
 - A set of generally accepted rules as to which registers will be unchanged after a procedure call (`jal`) and which may be changed.

Register Conventions

- **CalleR**: the calling function
- **CalleE**: the function being called
- When callee returns from executing, the caller needs to know which registers may have changed and which are guaranteed to be unchanged.

Caller vs. Callee

- **Caller has to save if it wants to preserve the value around a function call**
 - a0-a1 (can't put back in a0/a1 if return values use it)
 - a2-a7
 - t0-t6
 - ra: only needs to save once no mater how many calls it makes
- **Callee has to save if it wants to modify the value in its body**
 - s0 – s11
 - sp: handles this one differently – release what you allocate
 - gp: we will not use this, so not saved
 - tp: we will not use this, so not saved

Where to Save These Registers

- **The stack**
- **The code is in complete control of the stack pointer**
 - Allocate by
 - Release by
- **How does it know how much space it needs?**

How to Save sp?

- Procedure “knows” how much it changed sp by at the start
- Assumes every procedure it calls respects this also
- Procedure “restores” sp by restoring the change
- **Don’t store sp on the stack**

- If procedure changed sp and fails to restore sp, caller can’t locate its own stack items
- Recommend: only change sp at procedure entry and exit

RISC-V Registers

REGISTER NAME, USE, CALLING CONVENTION

| REGISTER | NAME | USE | SAVER |
|----------|----------|-------------------------------------|--------|
| x0 | zero | The constant value 0 | N.A. |
| x1 | ra | Return address | Caller |
| x2 | sp | Stack pointer | Callee |
| x3 | gp | Global pointer | -- |
| x4 | tp | Thread pointer | -- |
| x5-x7 | t0-t2 | Temporaries | Caller |
| x8 | s0/fp | Saved register/Frame pointer | Callee |
| x9 | s1 | Saved register | Callee |
| x10-x11 | a0-a1 | Function arguments/Return values | Caller |
| x12-x17 | a2-a7 | Function arguments | Caller |
| x18-x27 | s2-s11 | Saved registers | Callee |
| x28-x31 | t3-t6 | Temporaries | Caller |
| f0-f7 | ft0-ft7 | FP Temporaries | Caller |
| f8-f9 | fs0-fs1 | FP Saved registers | Callee |
| f10-f11 | fa0-fa1 | FP Function arguments/Return values | Caller |
| f12-f17 | fa2-fa7 | FP Function arguments | Caller |
| f18-f27 | fs2-fs11 | FP Saved registers | Callee |
| f28-f31 | ft8-ft11 | $R[rd] = R[rs1] + R[rs2]$ | Caller |

CSCI 341 Simplifications

- **Assume no one calls main**

- Main is only a caller, never a callee

- **gp, tp: don't worry about them**

- You can write perfectly good RISC-V code without them.

Registers- saved/preserved

- **s0-s11: Restore if callee changes.**
 - that's why they're called saved registers.
 - If the callee changes these in any way, it must restore the original values before returning.
- **sp: Restore if callee changes.**
 - Called must ensure it points to the same place before and after the jal call, so the caller can restore values from the stack.
- **HINT -- All saved registers start with s!**
- **Don't put sp on the stack**

Registers- volatile

- **ra**
 - The jal call itself will change this register.
 - Caller needs to save on stack if nested call.
- **a0-a1**
 - contain the new returned values.
- **a2-a7**
 - volatile argument registers.
 - Caller needs to save if need old values after the call.
- **t0-t6**
 - That's why they're called temporary
 - any procedure may change them at any time
 - Caller needs to save if they'll need same values after the call

Implications of Register Conventions

- If function R calls function E,
 - R must save any **temporary** registers that it may be using onto the stack before making a `jal` call.
 - E must save any S (**saved**) registers it intends to use before garbling up their values
- Remember: **Caller/callee need to save only temporary/saved registers they are using, not all registers.**

Steps for Making a Procedure Call (V.3, final version)

- 1) Caller saves needed volatile regs onto stack.**
- 2) Assign argument(s) in a0-a7, if any.**
- 3) jal ra, proc -- call procedure**
- 4) Callee allocates stack and preserves used s-reg**
- 5) Perform procedure's operations**
- 6) Callee releases stack and restores used s-reg**
- 7) Place result in a0, a1 if any, for caller**
- 8) jalr x0, ra, 0 – return to caller**
- 9) Caller restore values from stack.**

Procedure Call Conventions (Rules, to you)

- **Call with jal ra, proc; return with jalr x0, ra, 0**
 - Maintain the “fiction” of procedures with jal/jalr pairs
 - No other jumps between two procedure bodies, only within a procedure
- **Accept up to 8 arguments in a0-a7**
 - Always used in that order (can’t skip or get clever to avoid stack use)
 - We will never a 9th argument (simplifying assumption)
- **Return value is always in a0 (and a1 if two return values)**
- **Follow register conventions**
 - Even in function that only you call!
 - These dictate which registers will be unchanged after a procedure call and which may be changed
 - Assume the worst: that they are changed, every time!

Nested Procedure Call Example (C Code)

```
main() {  
    int i,j,k; /* i-k:s0-s2 */  
    ...  
    k = sumSquare(i,j); ...  
}  
  
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}  
  
int mult (int mcand, int mlier){  
    int product;  
    product = 0;  
    while (mlier > 0) {  
        product += mcand;  
        mlier -= 1; }  
    return product;  
}
```

RISC-V Code for main

main:

```
addi a0,s0,0          # arg0 = i
addi a1,s1,0          # arg1 = j
jal ra, sumSquare    # call sumSquare
addi s2,a0,0          # k = mult()
```

addi a7, x0, 10

ecall

```
main() {
int i,j,k,m; /* i-m:s0-s3 */
k = sumSquare(i,j);
}
```

Comments for main

- **main function ends with ecall 10, not jalr x0,ra,0, there is no need to save ra onto stack**
- **all variables used in main function are saved registers, so there's no need to save these onto stack**
- **main cannot rely on procedures not touching and changing any volatile registers – this includes a0-a7**

RISC-V Code for sumSquare

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y; }
```

sumSquare:

```
    addi sp,sp,-8    # space on stack  
    sw ra, 4(sp)      # save ret addr  
    sw a1, 0(sp)      # save y
```

```
    add a1,a0,zero # mult(x,x)  
    jal ra, mult      # call mult
```

```
    lw a1, 0(sp)      # restore y  
    add a0,a0,a1      # mult() +y  
    lw ra, 4(sp)      # get ret addr  
    addi sp,sp,8       # restore stack  
    jalr zero, ra, 0
```

sumSquare: Caller and Callee

- Has jal, so has to save/restore ra, and thus sp
- Only known values coming in are sp, ra, a0, a1
- No good choice for saved regs vs. temp regs., can't know what caller used vs. what callee used – so stick with our original conventions
 - s for local variables
 - t for intermediate results

RISC-V code for mult (leaf procedure)

mult:

addi t0,x0,0

prod=0

Loop:

bge 0,a1,Fin

if mlr <= 0, goto Fin

add t0,t0,a0

prod+=mc

addi a1,a1,-1

mlr-=1

jal zero,Loop

goto Loop

Fin:

addi a0,t0,0

a0=prod

jalr zero, ra, 0

return

```
int mult (int mcand, int mlier) {  
    int product = 0;  
    while (mlier > 0)  {  
        product += mcand;  
        mlier -= 1; }  
    return product;  
}
```

Comments for leaf procedure

- **mult is a leaf procedure!**
 - No jal calls are made from mult
- **No need to save/restore ra since it won't change**
- **Only known values coming in are ra, a0, a1**
- **Use volatile registers (t0-t6, also include a0-a7) for local variables in leaf procedures!**
- **if it uses saved registers, it must save/restore them – so avoid that, to avoid having to use the stack**
- **Why use temp registers not s registers for intermediate calculations?**
 - Since you are the compiler, bend the rules (t registers not just for intermediate results in a leaf)

Review Procedures So Far

- 1) Consider a nested procedure P that calls procedure Q. P should always save ra. **True or False?**
- 2) If P needs t0, and will write a0 to pass a parameter to Q, P might first
 - a) need to save a0 to the stack.
 - b) need to return to the main program
 - c) need to save t0 to the stack.
- 3) When P calls Q, P should expect
 - a) Q might pop less from the stack than Q pushed to the stack
 - b) Q will push to stack above the stack pointer
 - c) Q will pop from stack as much as Q pushed to the stack

Recursion: Caller IS Callee

- Each call creates a new stack frame
- special case of nested procedure call

Classic Factorial Problem in RISC-V

```
int fact (int n)
{
    if (n <= 1) return 1;
    else return (n * fact(n - 1));
}
```

- Argument n in a0
- Result in a0


| | | | |
|---------------------|-------------------------------------|-------|-------------------------|
| fact: | # Procedure Address 8000 | ← | |
| addi sp, sp, -8 | # adjust stack for 2 items | | Instructions 1 – 3 |
| sw ra, 4(sp) | # save the return address | | |
| sw a0, 0(sp) | # save the argument n | | |
| addi t0, zero, 1 | # t0 = 1 | | Instructions 4 – 5 |
| blt t0, a0, L1 | # if n > 1, go to L1 | | |
| addi a0, zero, 1 | # return 1 (n<=1) | | Instructions 6 – 8 |
| addi sp, sp, 8 | # restore stack pointer | | |
| jalr zero, ra, 0 | # return to caller | _____ | |
| L1: addi a0, a0, -1 | # n >= 1: argument gets (n - 1) | | Instructions 9 – 10 |
| jal ra, fact | # call fact with (n -1) | | |
| lw t0, 0(sp) | # return from jal: restore n to t0 | | |
| lw ra, 4(sp) | # restore the return address | | Instructions 11 – 13 |
| addi sp, sp, 8 | # adjust stack pointer: pop 2 items | | |
| mul a0, t0, a0 | # return n * fact (n - 1) | | Instructions 14 – 15 |
| jalr zero, ra, 0 | # return to the caller | | |

Main () calls fact () : First Call ($n = 3$)

- Assume `jal ra, fact` is located at location 3996 in `main()`
- Argument $n = 3$ in `a0`
- Return address for `main()` = _____ so `ra` = _____

- Instructions 1 – 3.
- Fill out the stack after the execution of the following instructions
- `sp` already adjusted in picture

```
addi sp, sp, -8      # adjust stack for 2 items  
sw    ra, 4(sp)       # save return address  
sw    a0, 0(sp)       # save argument
```



Instructions 4 – 5

- **Test for a0 > 1**
- **a0 = _____**

will blt jump or continue?

```
addi t0, zero, 1           # test for n > 1
blt  t0, a0, L1
```

Instructions 9 – 10

- **Jump to L1**
- **After executing instructions 9 and 10**

```
L1: addi a0, a0, -1      # n > 1: argument gets (n - 1)
    jal  ra, fact           # call fact with (n -1)
```


- a0 updated to _____²
- address of fact = _____8000
- ra = _____8040

fact() calls fact() : Second Call (n = 2)

- Argument n = 2 in a0
- Return address for fact() = _____, so ra = _____

- Instructions 1 – 3.
- Fill out the stack after the execution of the following instructions
- sp already adjusted

```
addi sp, sp, -8      # adjust stack for 2 items
sw    ra, 4(sp)       # save return address
sw    a0, 0(sp)       # save argument
```



Instructions 4 – 5

- Test for $a0 > 1$
 - $a0 = \underline{\hspace{2cm}}$
- will blt jump or continue?**

```
addi t0, zero, 1           # test for n > 1
blt  t0, a0, L1
```

Instructions 9 – 10

- **Jump to L1**
- **After executing instructions 9 and 10**

```
L1: addi a0,a0,-1          # n > 1: argument gets (n - 1)
    jal ra, fact            # call fact with (n -1)
```


- a0 updated to 1
- address of fact = 8000
- ra = 8040

fact() calls fact(): Third Call (n = 1)

- Argument n = 1 in a0
- Return address for fact() = _____ so (ra = _____)

- Instructions 1 – 3.
- Fill out the stack after the execution of the following instructions
- sp already adjusted

```
addi sp, sp, -8      # adjust stack for 2 items
sw    ra, 4(sp)       # save return address
sw    a0, 0(sp)       # save argument
```




Instructions 4 – 5

- Test for $a0 > 1$
 - $a0 = \underline{\hspace{2cm}}$
- will blt jump or continue?**

```
addi t0, zero, 1           # test for n > 1
blt  t0, a0, L1
```

Instructions 6 – 8

```
addi a0, zero, 1 # return 1  
addi sp, sp, 8   # restore stack pointer  
jalr zero, ra, 0 # return to caller
```




- Return to 8040 with a0 = 1
- New PC = 8040
- How many times are these 3 instructions executed over the duration of the program (until the completion of the calculation of the factorial)?
- We adjusted `sp` but never popped anything off. Why?

Instructions 11 – 13

```
lw t0, 0(sp)      # return from jal: recover argument n  
lw ra, 4(sp)      # restore the return address  
addi sp, sp, 8    # adjust stack pointer to pop 2 items
```

- **a0 =1 (from returned call)**
- **t0 = 2**
- **ra = 8040**
- **What is left in the stack after the execution of the instructions above?**



Instructions 14 - 15


```
mul a0, t0, a0          # return n * fact (n - 1)
jalr zero, ra, 0
```

- Return to fact() with a0 = 2
- New PC = 8040

Instructions 11 – 13

```
lw t0, 0(sp)      # return from jal: restore argument n  
lw ra, 4(sp)      # restore the return address  
addi sp, sp, 8    # adjust stack pointer to pop 2 items
```

- **a0 = 2 (from returned call)**
- **t0 = 3**
- **ra = 4000**
- **What is left in the stack after the execution of the instructions above?**



Instructions 14 - 15

```
mul a0, t0, a0    # return n * fact (n - 1)
jalr zero, ra, 0 # return to caller
```

- Return to main() with a0 = 6
- New PC = 4000

Recursive Procedure Call Example: Fibonacci Numbers

- The Fibonacci numbers are defined as follows:

$F(n) = F(n - 1) + F(n - 2),$
 $F(0)$ and $F(1)$ are defined to be 1

- Rewriting this in C :

```
int fib(int n) {  
    if (n == 0) { return 1; }  
    if (n == 1) { return 1; }  
    return (fib(n - 1) + fib(n - 2));  
}
```

Approaching the RISC-V Code

- **Decide what goes on the stack, write the prologue**
 - Allocating stack
 - Saving initial values
- **Write the epilogue**
 - Restore values from the stack
 - Release the stack
 - Return to the caller
- **Write the body—carefully**
 - This matches the C function body

Fibonacci Numbers (Prologue)

- Space for three words on the stack

- return addr (ra)
- n (a0)
- temp sum (s0)

- RISC-V code:

```
fib:  
    addi sp, sp, -12 # space for 3 words  
    sw ra, 8(sp)      # save return address  
    sw s0, 4(sp)      # save s0, used within
```

Fibonacci Numbers (Epilogue)

fib_done:

```
lw    s0, 4(sp)      # restore s0
lw    ra, 8(sp)      # restore ra
addi sp, sp, 12      # pop the stack frame
jalr zero, ra, 0    # return to caller
```

Fibonacci Numbers: Body Part 1

- Start with the base cases

```
if (n == 0) return 1;  
if (n == 1) return 1;
```

- RISC-V code:

```
addi t1, a0, 0          # save a0 in t1  
addi a0, x0, 1          # a0 = 1 (setup return value)  
beq t1, x0, fib_done   # if n==0 goto fib_done  
addi t0, x0, 1  
beq t1, t0, fib_done   # if n==1 goto fib_done
```

Fibonacci Numbers (Body Part 2)

- Next, the recursive case: careful, this is tricky

```
return fib(n-1) + fib(n-2)
```

- RISC-V code:

```
addi a0, t1, -1      # a0 = n-1
sw   a0, 0(sp)       # need a0 after jal
jal  ra, fib         # fib(n-1)
add  s0, a0, 0        # save result
lw   a0, 0(sp)       # restore a0
addi a0, a0, -1      # a0 = n-2 (already n-1)
jal  ra, fib         # fib(n-2)
add  a0, a0, s0       # a0=fib(n-2)+fib(n-1)
                      # continues at fib_done...
```

- **Breakpoint on jal's and jalr**
- **Watch a0, s0, sp, ra**
- **Single step the first jal to see ra change (on first call)**
 - Updated, but stays unchanged, on the recursive call
- **Single step the jalr x0, ra, 0 to see PC change (on last return)**
 - Updated, but stays unchanged, on the recursive returns

RISC-V Register Conventions

| Name | Register number | Usage | Preserved on call? |
|-----------|-----------------|--------------------------------|--------------------|
| x0 | 0 | The constant value 0 | n.a. |
| x1 (ra) | 1 | Return address (link register) | yes |
| x2 (sp) | 2 | Stack pointer | yes |
| x3 (gp) | 3 | Global pointer | yes |
| x4 (tp) | 4 | Thread pointer | yes |
| x5 - x7 | 5–7 | Temporaries | no |
| x8 - x9 | 8–9 | Saved | yes |
| x10 - x17 | 10–17 | Arguments/results | no |
| x18 - x27 | 18–27 | Saved | yes |
| x28 - x31 | 28–31 | Temporaries | no |

After-Class Activities

- Draw out the stack showing the execution of fib.s with an input of 4 (the first 4 Fibonacci numbers will be printed)
- Compare with what RARS stack shows